

A Comprehensive Field Guide for Data Scientists

February 2026

Contents

Preface	3
1 Cross-Validation: Honest Performance Estimation	4
1.1 The Core Problem: Why a Single Train–Test Split Is Not Enough	4
1.2 K-Fold Cross-Validation	4
1.2.1 How It Works	4
1.2.2 Choosing k	4
1.2.3 Python Example: K-Fold and Stratified K-Fold	5
1.3.1 Walk-Forward Validation (Expanding Window)	7
2 Hyperparameter Tuning: Systematic Model Configuration	10
2.1 What Are Hyperparameters?	10
2.2 Grid Search	10
2.3 Random Search	12
2.4 Bayesian Optimisation	13
2.4.1 Choosing a Tuning Strategy	15
3 Metrics for Imbalanced Datasets	17
3.1 Why Accuracy Fails	17
3.2 The Confusion Matrix: A Foundation	17
3.3 Core Metrics	17
3.3.1 Reading Precision-Recall Curves	21
3.3.2 Threshold Optimisation	21
4 Model Interpretability	22
4.1 Why Interpretability Is Not Optional	22
4.2 SHAP: SHapley Additive exPlanations	22
4.2.1 The Theory	22
4.2.2 Reading SHAP Output	25
4.3 LIME: Local Interpretable Model-agnostic Explanations	25
4.3.1 SHAP vs LIME: When to Use Each	28
4.4 Partial Dependence Plots	28
5 Putting It All Together: End-to-End Evaluation Workflow	32
6 The Evaluation Checklist	36
7 Regression Metrics: Measuring Continuous Predictions	38
7.1 The Problem with a Single Number	38
7.2 Core Regression Metrics	38
7.3 Residual Analysis: Beyond a Single Number	39
8 Model Calibration: When Probabilities Must Mean Something	44
8.1 The Calibration Problem	44
8.2 The Reliability Diagram	44
8.3 The Brier Score	44

9 Statistical Significance Testing for Model Comparison	48
9.1 The Core Question: Is Model A Actually Better?	48
9.2 McNemar's Test: Comparing Classifiers on the Same Test Set	48
9.3 The Corrected Resampled t-Test	48
10 Bias & Fairness Evaluation	52
10.1 Why Fairness Is a Measurement Problem	52
10.2 Core Fairness Definitions	52
11 Production Monitoring & Drift Detection	57
11.1 The Life Cycle Does Not End at Deployment	57
11.2 Population Stability Index (PSI)	57
12 Learning Curves & Diagnosing Underfitting vs Overfitting	62
12.1 The Bias-Variance Trade-off in Practice	62
12.2 Reading Learning Curves	62
12.3 The Validation Curve: Finding the Complexity Sweet Spot	65
13 Multiclass & Multi-label Metrics	66
13.1 Beyond Binary Classification	66
13.2 Averaging Strategies	66
13.3 Cohen's Kappa: Agreement Beyond Chance	66
13.4 The Normalised Confusion Matrix as a Debugging Tool	69
Appendix: Quick Reference	72

Preface: Why Evaluation Is the Most Important Skill

Building a machine learning model that produces numbers is easy. Building one you can *trust*, defend in a board meeting, and rely on to drive real decisions is an entirely different discipline. That discipline is **Model Evaluation & Validation**.

Many data science projects fail not because the algorithm was wrong, but because the practitioner optimised for the wrong metric, let data from the future leak into training, reported accuracy on a wildly imbalanced test set, or could not explain to a stakeholder why the model made a particular decision. This tutorial is designed to close those gaps.

The guide is structured around four pillars, each building on the previous one:

1. **Cross-Validation** — learning how to estimate generalisation performance honestly, and how to avoid the most common pitfalls like data leakage and overfitting to a validation fold.
2. **Hyperparameter Tuning** — systematically searching the space of model configurations rather than guessing, using methods that scale from exhaustive search to Bayesian optimisation.
3. **Metrics for Imbalanced Datasets** — understanding why accuracy is misleading when classes are unequal and which metrics actually reflect business cost.
4. **Model Interpretability** — opening the black box with SHAP, LIME, and partial dependence plots so that the model can be scrutinised, debugged, and explained.

Throughout, every concept is grounded in runnable Python examples using mainstream libraries (`scikit-learn`, `shap`, `lime`, `scikit-optimize`). Code is annotated line by line wherever a concept needs emphasis.

1 Cross-Validation: Honest Performance Estimation

1.1 The Core Problem: Why a Single Train–Test Split Is Not Enough

Suppose you train a classifier and measure its accuracy on a held-out test set. You get 87%. Should you celebrate? Not yet. That number reflects one particular random partition of your data. Had you used a different random seed, you might have gotten 82% or 91%. You have measured the performance of *one* model on *one* subset of your data, not the expected performance of your modelling *approach* on unseen data.

Cross-validation resolves this by evaluating the model on multiple, non-overlapping validation subsets and averaging the results. The variance of those scores tells you how stable your model is; the mean tells you the unbiased estimate of generalisation performance.

Just as a model can overfit the training data, an *evaluation procedure* can overfit the validation data. Every time you make a modelling decision based on a fixed validation set, you implicitly “train” on that set. Cross-validation delays this leakage by rotating which data is held out.

1.2 K-Fold Cross-Validation

1.2.1 How It Works

K-Fold partitions the dataset into k equally sized **folds**. The model is trained k times; each time, one fold serves as the validation set and the remaining $k - 1$ folds form the training set. The final performance estimate is the mean (and standard deviation) across all k scores.

Fold 1	VAL	TRAIN	TRAIN	TRAIN	TRAIN
Fold 2	TRAIN	VAL	TRAIN	TRAIN	TRAIN
Fold 3	TRAIN	TRAIN	VAL	TRAIN	TRAIN
Fold 4	TRAIN	TRAIN	TRAIN	VAL	TRAIN
Fold 5	⇒ TRAIN	TRAIN	TRAIN	TRAIN	VAL

Figure 1: 5-Fold Cross-Validation: the teal fold rotates through each iteration.

1.2.2 Choosing k

- **$k = 5$ or 10** is the standard choice backed by empirical research (Kohavi 1995). It balances bias and variance in the estimator.
- **$k = n$** (Leave-One-Out CV) is nearly unbiased but has very high variance and is computationally expensive.
- **Stratified K-Fold** preserves the class distribution in each fold — always use this for classification.

1.2.3 Python Example: K-Fold and Stratified K-Fold

K-Fold Cross-Validation — Full Walkthrough

```

1 import numpy as np
2 import pandas as pd
3 from sklearn.datasets import make_classification
4 from sklearn.ensemble import RandomForestClassifier
5 from sklearn.model_selection import (
6     KFold, StratifiedKFold, cross_val_score, cross_validate
7 )
8 from sklearn.preprocessing import StandardScaler
9 from sklearn.pipeline import Pipeline
10
11 # -----
12 # 1. Synthetic dataset: 1000 samples, 20 features, binary target
13 #     weights=(0.7, 0.3) means 70% class 0, 30% class 1 mild
14 #     imbalance
15 X, y = make_classification(
16     n_samples=1000,
17     n_features=20,
18     n_informative=10,
19     n_redundant=4,
20     weights=[0.7, 0.3],
21     random_state=42
22 )
23
24 # -----
25 # 2. Build a Pipeline
26 #     Always wrap preprocessing inside the pipeline so that the
27 #     scaler
28 #         is fit ONLY on training folds never on the validation fold.
29 #         This is the single most common source of data leakage.
30 # -----
31 pipeline = Pipeline([
32     ('scaler', StandardScaler()),                      # zero-mean, unit-
33     ('clf', RandomForestClassifier(                  variance
34         n_estimators=100, random_state=42
35     ))
36 ]
37 )
38 # -----
39 # 3. Plain K-Fold (ignores class distribution in each fold)
40 kf = KFold(n_splits=5, shuffle=True, random_state=42)
41
42 scores_kfold = cross_val_score(
43     pipeline, X, y,
44     cv=kf,
45     scoring='roc_auc',      # area under ROC curve: threshold-
46     independent
47     n_jobs=-1             # use all CPU cores
48 )
49 print("== Standard K-Fold ==")
50 print(f"Fold AUC scores : {np.round(scores_kfold, 4)}")

```

```

51  print(f"Mean AUC           : {scores_kfold.mean():.4f}")
52  print(f"Std  AUC            : {scores_kfold.std():.4f}")
53
54  # -----
55  # 4. Stratified K-Fold (preserves class ratio in every fold)
56  #      For classification, ALWAYS prefer StratifiedKFold.
57  #
58  skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
59
60  scores_stratified = cross_val_score(
61      pipeline, X, y,
62      cv=skf,
63      scoring='roc_auc',
64      n_jobs=-1
65  )
66
67  print("\n==== Stratified K-Fold ====")
68  print(f"Fold AUC scores : {np.round(scores_stratified, 4)}")
69  print(f"Mean AUC         : {scores_stratified.mean():.4f}")
70  print(f"Std  AUC          : {scores_stratified.std():.4f}")
71
72  # -----
73  # 5. cross_validate returns richer information:
74  #      train scores, test scores, and fit/score times.
75  #
76  cv_results = cross_validate(
77      pipeline, X, y,
78      cv=skf,
79      scoring=['roc_auc', 'f1', 'precision', 'recall'],
80      return_train_score=True,    # detect overfitting by comparing
81      n_jobs=-1
82  )
83
84  results_df = pd.DataFrame({
85      'train_roc_auc': cv_results['train_roc_auc'],
86      'test_roc_auc' : cv_results['test_roc_auc'],
87      'test_f1'       : cv_results['test_f1'],
88      'fit_time_s'   : cv_results['fit_time']
89  })
90  print("\n==== Detailed CV Results ====")
91  print(results_df.round(4).to_string(index=False))
92
93  # If train AUC >> test AUC, the model is overfitting.
94  gap = cv_results['train_roc_auc'].mean() - cv_results['test_roc_auc'][
95  ].mean()
96  print(f"\nOverfit gap (train - test AUC): {gap:.4f}")
97  if gap > 0.05:
98      print("WARNING: Possible overfitting detected.")

```

What this code does, step by step:

Step 1 creates a synthetic binary classification dataset with mild class imbalance, which mirrors real business scenarios where one outcome (e.g., fraud, churn) is rarer than the other.

Step 2 is the most important architectural decision: the `Pipeline` ensures that the `StandardScaler` is fitted *only* on the training portion of each fold. If you scaled all data

first and then cross-validated, the validation fold's statistics would have influenced the scaler — a subtle form of data leakage that inflates your score.

Steps 3 and 4 compare plain `KFold` with `StratifiedKFold`. In practice, stratification reduces the variance of the CV estimator when class imbalance is present.

Step 5 uses

1.3 Time-Series Cross-Validation

Standard K-Fold **must not be used with time-series data**. The reason is fundamental: the future cannot cause the past. If your validation fold contains data from January and your training fold contains data from March, your model learns from the future to predict the past. This produces wildly optimistic estimates that will not hold in live deployment.

1.3.1 Walk-Forward Validation (Expanding Window)

`TimeSeriesSplit` implements ***walk-forward validation***, also called *expanding window* or *rolling origin* cross-validation. The training window always ends before the validation window begins, and the training set grows with each split.

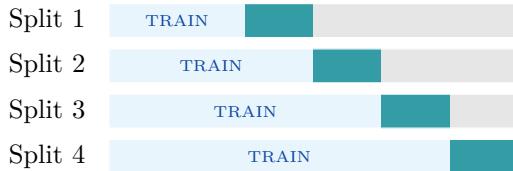


Figure 2: Time-Series Walk-Forward Validation: training always precedes validation temporally.

Time-Series Cross-Validation

```

1 import numpy as np
2 import pandas as pd
3 from sklearn.model_selection import TimeSeriesSplit, cross_validate
4 from sklearn.ensemble import GradientBoostingClassifier
5 from sklearn.pipeline import Pipeline
6 from sklearn.preprocessing import StandardScaler
7
8 # -----
9 # 1. Simulate a time-ordered dataset (e.g., daily transaction data)
10 #     n_samples = 2000 days; features represent behavioural signals.
11 # -----
12 np.random.seed(42)
13 n = 2000
14
15 # Create a temporal trend in the signal to mimic real time-series
16 time_idx = np.arange(n)
17 signal = 0.3 * np.sin(2 * np.pi * time_idx / 365) # seasonal
18         component
19 X = np.column_stack([
20     signal + np.random.randn(n) * 0.5,    # feature 1: noisy signal
21     np.random.randn(n),                  # feature 2: noise

```

```

22         time_idx / n + np.random.randn(n)*0.1 # feature 3: trend
23     ])
24 y = (signal + np.random.randn(n) * 0.4 > 0).astype(int)
25
26 # -----
27 # 2. Configure TimeSeriesSplit
28 #     n_splits=5 creates 5 folds, each with progressively more
29 #     training
30 #     data. gap=30 inserts a 30-day gap between training and
31 #     validation
32 #     to simulate realistic deployment delay.
33 # -----
34 tscv = TimeSeriesSplit(n_splits=5, gap=30)
35
36 # -----
37 # 3. Inspect what each split looks like
38 # -----
39 print("Split summary (train size | val size):")
40 for fold_num, (train_idx, val_idx) in enumerate(tscv.split(X), start
41 =1):
42     print(f" Fold {fold_num}: train [{train_idx[0]}..{train_idx[-1]}
43         }] "
44         f"({len(train_idx)} samples) | "
45         f"val [{val_idx[0]}..{val_idx[-1]}] ({len(val_idx)}
46             samples)")
47
48 # -----
49 # 4. Cross-validate a gradient boosting classifier
50 # -----
51 pipeline = Pipeline([
52     ('scaler', StandardScaler()),
53     ('clf', GradientBoostingClassifier(
54         n_estimators=100, learning_rate=0.05,
55         max_depth=4, random_state=42
56     )))
57 ])
58
59 cv_results = cross_validate(
60     pipeline, X, y,
61     cv=tscv,
62     scoring=['roc_auc', 'f1'],
63     return_train_score=True
64 )
65
66 print("\n==== Time-Series CV Results ====")
67 for fold_i in range(5):
68     print(f" Fold {fold_i+1}: "
69         f"train AUC={cv_results['train_roc_auc'][fold_i]:.4f} "
70         f"val AUC={cv_results['test_roc_auc'][fold_i]:.4f} "
71         f"val F1={cv_results['test_f1'][fold_i]:.4f}")
72
73 print(f"\nMean val AUC : {cv_results['test_roc_auc'].mean():.4f} "
74       f"+/- {cv_results['test_roc_auc'].std():.4f}")
75
76 # -----
77 # 5. Trend analysis: does performance degrade over time?
78 #     In live systems, concept drift causes later folds to score
79 #     lower.

```

```
74 # -----
75 fold_aucs = cv_results['test_roc_auc']
76 trend = np.polyfit(range(len(fold_aucs)), fold_aucs, deg=1)[0]
77 print(f"\nPerformance trend across folds: {trend:+.4f} per fold")
78 if trend < -0.01:
79     print("ALERT: Declining performance potential concept drift.")
80 else:
81     print("Performance stable across time folds.")
```

The gap parameter is critical for production systems. If your model is retrained monthly and predictions are served for the following 30 days, inserting a 30-day gap into your CV scheme ensures that your offline evaluation mimics live conditions. Without it, validation samples immediately adjacent to the training cutoff may share patterns (such as the same weekly seasonality cycle) that won't exist at actual deployment time.

Trend analysis across folds is a uniquely time-series concern. If AUC falls monotonically from fold 1 to fold 5, your model is experiencing **concept drift** — the statistical relationship between features and target is changing over time. This is a signal to explore re-training schedules or online learning approaches.

- **Shuffling data** before splitting — destroys temporal ordering and guarantees leakage.
- **Using aggregated features** (e.g., 30-day rolling mean computed on the full dataset) before splitting. Always compute such features inside the pipeline or after splitting.
- **Ignoring the deployment gap** between training cutoff and live prediction.

2 Hyperparameter Tuning: Systematic Model Configuration

2.1 What Are Hyperparameters?

Machine learning models have two kinds of parameters. **Model parameters** are learned during training (e.g., the weights in a neural network, the split thresholds in a decision tree). **Hyperparameters** are set *before* training begins and control the learning process itself: how many trees to grow, how deep each tree can be, the learning rate, regularisation strength, and so on.

The goal of hyperparameter tuning is to find the configuration that maximises the model's generalisation performance, as measured by cross-validation, not by training error.

Every tuning method follows the same loop: **propose** a configuration → **train** the model with that configuration → **evaluate** using cross-validation → **record** the score → repeat until budget exhausted → **select** the configuration with the best CV score. The methods differ in how they *propose* the next candidate.

2.2 Grid Search

Grid search exhaustively evaluates every combination of a pre-specified set of hyperparameter values. If you specify 3 values for `max_depth` and 4 values for `n_estimators`, grid search trains $3 \times 4 = 12$ models (each cross-validated k times), giving $12k$ total fits.

This is guaranteed to find the best configuration *within the grid*, but it scales exponentially with the number of hyperparameters – a phenomenon called the **curse of dimensionality in search**.

Grid Search with Cross-Validation

```

1 import numpy as np
2 import pandas as pd
3 from sklearn.datasets import make_classification
4 from sklearn.ensemble import GradientBoostingClassifier
5 from sklearn.model_selection import GridSearchCV, StratifiedKFold
6 from sklearn.pipeline import Pipeline
7 from sklearn.preprocessing import StandardScaler
8
9 X, y = make_classification(
10     n_samples=1500, n_features=20, n_informative=12,
11     weights=[0.65, 0.35], random_state=42
12 )
13
14 # -----
15 # Pipeline: scaler + classifier
16 # Hyperparameters for pipeline steps are referenced as
17 # "stepname__parametername"
18 # -----
19 pipeline = Pipeline([

```

```

20     ('scaler', StandardScaler()),
21     ('clf', GradientBoostingClassifier(random_state=42))
22 )
23
24 # -----
25 # Define the grid
26 # 3 x 4 x 3 = 36 configurations, each with 5-fold CV = 180 fits.
27 # For production models, this grid is modest; real searches often
28 # involve hundreds of configurations.
29 #
30 param_grid = {
31     'clf__n_estimators': [50, 100, 200],
32     'clf__max_depth': [2, 3, 4, 5],
33     'clf__learning_rate': [0.01, 0.05, 0.1],
34 }
35
36 cv_strategy = StratifiedKFold(n_splits=5, shuffle=True, random_state
=42)
37
38 grid_search = GridSearchCV(
39     estimator=pipeline,
40     param_grid=param_grid,
41     scoring='roc_auc',           # optimise for AUC
42     cv=cv_strategy,
43     n_jobs=-1,                 # parallelise across all cores
44     verbose=1,                  # print progress
45     refit=True                  # refit on full data with best params
46 )
47
48 grid_search.fit(X, y)
49
50 # -----
51 # Results exploration
52 #
53 print(f"Best CV AUC : {grid_search.best_score_:.4f}")
54 print(f"Best params : {grid_search.best_params_}")
55
56 # Convert the full results to a DataFrame for analysis
57 results_df = pd.DataFrame(grid_search.cv_results_)
58
59 # Show top-5 configurations
60 top5 = results_df[['params', 'mean_test_score', 'std_test_score',
61                   'rank_test_score']]\
62     .sort_values('rank_test_score').head(5)
63
64 print("\nTop 5 configurations:")
65 for _, row in top5.iterrows():
66     print(f" Rank {int(row['rank_test_score'])}: "
67           f"AUC={row['mean_test_score']:.4f} "
68           f"+/-{row['std_test_score']:.4f} "
69           f"{row['params']} ")
70
71 # -----
72 # The best estimator is already refitted on the full dataset.
73 # It can be used directly for prediction.
74 #
75 best_model = grid_search.best_estimator_

```

2.3 Random Search

Random search samples hyperparameter combinations uniformly at random from a defined distribution. Bergstra and Bengio (2012) demonstrated mathematically that random search is more efficient than grid search when the number of hyperparameters is large and only a few are truly important, because random search explores more of the *marginal* space of each parameter with the same number of evaluations.

The intuition: if 2 out of 5 hyperparameters matter but you don't know which 2, a $5 \times 5 \times 5 \times 5 \times 5$ grid evaluates each important parameter at only 5 distinct values. Fifty random samples explore each important parameter at up to 50 distinct values, while still covering the full joint space.

Random Search with Continuous Distributions

```

1  import numpy as np
2  from scipy.stats import uniform, randint, loguniform
3  from sklearn.ensemble import GradientBoostingClassifier
4  from sklearn.model_selection import RandomizedSearchCV,
5      StratifiedKFold
6  from sklearn.pipeline import Pipeline
7  from sklearn.preprocessing import StandardScaler
8  from sklearn.datasets import make_classification
9
10 X, y = make_classification(
11     n_samples=1500, n_features=20, n_informative=12,
12     weights=[0.65, 0.35], random_state=42
13 )
14
15 pipeline = Pipeline([
16     ('scaler', StandardScaler()),
17     ('clf', GradientBoostingClassifier(random_state=42))
18 ])
19 # -----
20 # Distributions instead of grids
21 #
22 # loguniform(a, b): draws from a log-uniform distribution between a
23 # and b, which is appropriate for learning rates and
24 # regularisation
25 # strengths that span orders of magnitude.
26 #
27 # randint(low, high): draws integers from [low, high].
28 # Best for counts like n_estimators, max_depth.
29 #
30 # uniform(loc, scale): draws from [loc, loc+scale].
31 # Best for continuous values with no strong scale preference.
32 #
33 param_distributions = {
34     'clf__n_estimators': randint(50, 500),           # integers in [
35             50, 500)
36     'clf__max_depth': randint(2, 8),                 # integers in [
37             2, 8)
38     'clf__learning_rate': loguniform(1e-3, 0.5),    # log-scale [0.
39             001, 0.5]
40     'clf__subsample': uniform(0.5, 0.5),            # uniform in [0
41             .5, 1.0]

```

```

37     'clf__min_samples_leaf': randint(1, 30),
38     'clf__max_features': uniform(0.3, 0.7),           # fraction of
39   }
40
41 cv_strategy = StratifiedKFold(n_splits=5, shuffle=True, random_state
42 =42)
43
44 random_search = RandomizedSearchCV(
45     estimator=pipeline,
46     param_distributions=param_distributions,
47     n_iter=60,           # evaluate 60 random configurations
48     scoring='roc_auc',
49     cv=cv_strategy,
50     n_jobs=-1,
51     random_state=42,
52     verbose=1,
53     refit=True
54 )
55
56 random_search.fit(X, y)
57
58 print(f"Best CV AUC : {random_search.best_score_:.4f}")
59 print(f"Best params :")
60 for k, v in random_search.best_params_.items():
61     print(f"  {k}: {v}")
62
63 # -----
64 # Compare grid search vs random search coverage
65 # For 60 evaluations:
66 #   Grid: constrained to pre-specified values
67 #   Random: explores continuous ranges, likely finds better optima
68 # -----

```

2.4 Bayesian Optimisation

Bayesian optimisation is the state of the art for expensive hyperparameter tuning. Unlike grid and random search, which treat each evaluation independently, Bayesian optimisation *learns* from previous evaluations and uses a probabilistic **surrogate model** (typically a Gaussian Process or Tree-structured Parzen Estimator) to predict which regions of the hyperparameter space are most promising.

The surrogate model balances **exploitation** (evaluating where the model predicts the best performance) with **exploration** (evaluating uncertain regions that might harbour better optima). This trade-off is managed by an **acquisition function**, most commonly Expected Improvement (EI):

$$\text{EI}(\mathbf{x}) = \mathbb{E}[\max(f(\mathbf{x}) - f(\mathbf{x}^+), 0)]$$

where $f(\mathbf{x}^+)$ is the best observed score so far. The next evaluation point is $\mathbf{x}_{\text{next}} = \arg \max_{\mathbf{x}} \text{EI}(\mathbf{x})$.

Bayesian Optimisation with scikit-optimize

```

1 import numpy as np
2 from sklearn.datasets import make_classification
3 from sklearn.ensemble import GradientBoostingClassifier
4 from sklearn.model_selection import StratifiedKFold, cross_val_score
5 from sklearn.pipeline import Pipeline
6 from sklearn.preprocessing import StandardScaler
7
8 # scikit-optimize provides BayesSearchCV with a scikit-learn API
9 # Install: pip install scikit-optimize
10 from skopt import BayesSearchCV
11 from skopt.space import Real, Integer, Categorical
12
13 X, y = make_classification(
14     n_samples=1500, n_features=20, n_informative=12,
15     weights=[0.65, 0.35], random_state=42
16 )
17
18 pipeline = Pipeline([
19     ('scaler', StandardScaler()),
20     ('clf', GradientBoostingClassifier(random_state=42))
21 ])
22
23 # -----
24 # Search space using skopt types
25 #   Real(low, high, prior='log-uniform'): continuous log-scale range
26 #   Integer(low, high): integer range
27 #   Categorical([...]): discrete set of choices
28 #
29 search_space = {
30     'clf__n_estimators' : Integer(50, 400),
31     'clf__max_depth'   : Integer(2, 7),
32     'clf__learning_rate': Real(1e-3, 0.5, prior='log-uniform'),
33     'clf__subsample'    : Real(0.4, 1.0),
34     'clf__min_samples_leaf': Integer(1, 25),
35 }
36
37 cv_strategy = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
38
39 # -----
40 # BayesSearchCV: n_iter controls total evaluations.
41 # Fewer evaluations than grid/random search can achieve better
42 # results
43 # because each evaluation is informed by prior results.
44 bayes_search = BayesSearchCV(
45     estimator=pipeline,
46     search_spaces=search_space,
47     n_iter=40,                      # 40 informed evaluations
48     scoring='roc_auc',
49     cv=cv_strategy,
50     n_jobs=1,                      # skopt internals require n_jobs=1
51     random_state=42,
52     verbose=0,
53     refit=True
54 )

```

```

55
56 bayes_search.fit(X, y)
57
58 print(f"Best Bayesian CV AUC : {bayes_search.best_score_:.4f}")
59 print("Best parameters:")
60 for k, v in bayes_search.best_params_.items():
61     print(f"  {k}: {v}")
62
63 # -----
64 # Convergence plot: track how AUC improves iteration by iteration.
65 # A flat curve means the search has converged; more budget won't
66 # help.
67 # -----
68 import matplotlib.pyplot as plt
69
70 scores_over_time = []
71 best_so_far = -np.inf
72 for score in bayes_search.cv_results_['mean_test_score']:
73     best_so_far = max(best_so_far, score)
74     scores_over_time.append(best_so_far)
75
76 plt.figure(figsize=(8, 4))
77 plt.plot(range(1, len(scores_over_time)+1), scores_over_time,
78          marker='o', color='#0d47a1', linewidth=2, markersize=4)
79 plt.xlabel("Evaluation number")
80 plt.ylabel("Best CV AUC so far")
81 plt.title("Bayesian Optimisation Convergence")
82 plt.grid(alpha=0.3)
83 plt.tight_layout()
84 plt.savefig("bayesian_convergence.png", dpi=150)
85 plt.close()
86 print("Convergence plot saved.")

```

2.4.1 Choosing a Tuning Strategy

The following table summarises when to use each approach:

Method	Best for	Weakness	Budget
Grid Search	Small grids, interpretable search space	Exponential scaling	Low
Random Search	Many hyperparameters, few matter	May miss fine structure	Medium
Bayesian Opt.	Expensive models, large spaces	Sequential (harder to parallelise)	High

Table 1: Hyperparameter tuning method comparison.

When you tune hyperparameters with 5-fold stratified CV, you must also report your final performance estimate using a separate *outer* cross-validation loop (nested CV), or a hold-out test set that was never touched during tuning. Using the same data for both tuning and final reporting is a form of overfitting.

3 Metrics for Imbalanced Datasets

3.1 Why Accuracy Fails

Imagine a fraud detection model. Fraudulent transactions represent 0.5% of all transactions. A model that predicts “not fraud” for every single transaction achieves 99.5% accuracy. This is a useless model that catches zero fraud.

This is why **class imbalance** is one of the most practically important challenges in applied machine learning. The choice of evaluation metric must reflect the cost structure of the problem, not just raw prediction accuracy.

3.2 The Confusion Matrix: A Foundation

All classification metrics derive from four counts:

	Predicted Positive	Predicted Negative
Actual Positive	green!15 TP (True Positive)	red!15 FN (False Negative)
Actual Negative	red!15 FP (False Positive)	green!15 TN (True Negative)

Table 2: The 2×2 confusion matrix.

From these four numbers, we derive every metric in this section.

3.3 Core Metrics

Precision answers the question: “Of all the cases I predicted positive, what fraction actually were positive?” It penalises false alarms.

$$\text{Precision} = \frac{TP}{TP + FP}$$

Recall (also called Sensitivity or True Positive Rate) answers: “Of all the actual positive cases, what fraction did I catch?” It penalises missed detections.

$$\text{Recall} = \frac{TP}{TP + FN}$$

F1-Score is the harmonic mean of precision and recall. The harmonic mean is used (rather than arithmetic) because it punishes extreme imbalances between precision and recall.

$$F_1 = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$$

The F_β score generalises F1 by weighting recall β times more than precision. For fraud detection where missing a fraud (false negative) is far more costly than a false alarm, use $\beta > 1$:

$$F_\beta = (1 + \beta^2) \cdot \frac{\text{Precision} \cdot \text{Recall}}{\beta^2 \cdot \text{Precision} + \text{Recall}}$$

Balanced Accuracy averages recall across all classes, correcting for imbalance without requiring threshold tuning:

$$\text{Balanced Accuracy} = \frac{1}{2} \left(\frac{TP}{TP + FN} + \frac{TN}{TN + FP} \right)$$

Comprehensive Imbalanced Classification Evaluation

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3 from sklearn.datasets import make_classification
4 from sklearn.ensemble import GradientBoostingClassifier,
    RandomForestClassifier
5 from sklearn.linear_model import LogisticRegression
6 from sklearn.model_selection import StratifiedKFold, cross_val_
    predict
7 from sklearn.pipeline import Pipeline
8 from sklearn.preprocessing import StandardScaler
9 from sklearn.metrics import (
    classification_report,
    confusion_matrix,
    f1_score,
    balanced_accuracy_score,
    precision_recall_curve,
    average_precision_score,
    roc_auc_score,
    ConfusionMatrixDisplay
)
19
20 # -----
21 # 1. Heavily imbalanced dataset: 95% class 0, 5% class 1
22 #     This simulates fraud, disease, or rare equipment failure.
23 #
24 X, y = make_classification(
25     n_samples=5000,
26     n_features=20,
27     n_informative=10,
28     weights=[0.95, 0.05],    # 4750 negatives, 250 positives
29     random_state=42
30 )
31 print(f"Class distribution: {np.bincount(y)}  "
32       f"({100*y.mean():.1f}% positive)")
33
34 # -----
35 # 2. Train three models and get out-of-fold predictions
36 #     cross_val_predict generates predictions for every sample using
37 #     only models trained without that sample proper OOF evaluation.
38 #

```

```

39  models = {
40      'Logistic Regression': Pipeline([
41          ('scaler', StandardScaler()),
42          ('clf', LogisticRegression(
43              class_weight='balanced', # upweight minority
44              C=0.1, max_iter=1000
45          ))
46      ],
47      'Random Forest': Pipeline([
48          ('scaler', StandardScaler()),
49          ('clf', RandomForestClassifier(
50              n_estimators=200,
51              class_weight='balanced_subsample',
52              random_state=42
53          ))
54      ],
55      'Gradient Boosting': Pipeline([
56          ('scaler', StandardScaler()),
57          ('clf', GradientBoostingClassifier(
58              n_estimators=200, learning_rate=0.05,
59              max_depth=4, random_state=42
60          ))
61      ],
62  })
63
64  cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
65
66  results = {}
67  for name, pipe in models.items():
68      # predict_proba: probability scores for each class
69      proba = cross_val_predict(pipe, X, y, cv=cv, method='predict_'
70          proba')[ :, 1]
71      # predict: hard class labels at 0.5 threshold
72      labels = cross_val_predict(pipe, X, y, cv=cv, method='predict')
73
74      results[name] = {
75          'proba' : proba,
76          'labels' : labels,
77          'ap' : average_precision_score(y, proba), # area under
78          'P-R curve',
79          'auc' : roc_auc_score(y, proba),
80          'f1' : f1_score(y, labels),
81          'f2' : f1_score(y, labels, beta=2.0), # recall-
82          'weighted',
83          'bal_acc' : balanced_accuracy_score(y, labels),
84      }
85
86  # -----
87  # 3. Summary table
88  # -----
89  print("\n{:<22} {:>8} {:>8} {:>8} {:>8} {:>8} ".format(
90      "Model", "AUC", "Avg Prec", "F1", "F2", "BalAcc"))
91  print("-" * 70)
92  for name, r in results.items():
93      print(f"{{name:<22} {{r['auc']:>8.4f} {{r['ap']:>8.4f} "
94          f"{{r['f1']:>8.4f} {{r['f2']:>8.4f} {{r['bal_acc']:>8.4f}}}")
95
96  # -----

```

```

94 # 4. Precision-Recall curves the right plot for imbalanced data
95 # ROC curves are optimistic when the negative class dominates;
96 # P-R curves expose the real difficulty of finding true positives
97 .
98 # -----
99 fig, axes = plt.subplots(1, 2, figsize=(13, 5))
100
101 colors = ['#0d47a1', '#00838f', '#e65100']
102 for (name, r), color in zip(results.items(), colors):
103     prec, rec, thresh = precision_recall_curve(y, r['proba'])
104     axes[0].plot(rec, prec, color=color, linewidth=2,
105                  label=f'{name} (AP={r["ap"]:.3f})')
106
107 axes[0].axhline(y.mean(), color='gray', linestyle='--', linewidth=1.
108                  2,
109                  label=f"Random (AP={y.mean():.3f})")
110 axes[0].set_xlabel("Recall", fontsize=12)
111 axes[0].set_ylabel("Precision", fontsize=12)
112 axes[0].set_title("Precision-Recall Curves", fontsize=13, fontweight
113                  = 'bold')
114 axes[0].legend(fontsize=9)
115 axes[0].grid(alpha=0.3)
116
117 # -----
118 # 5. Threshold analysis for the best model
119 # The default 0.5 threshold is rarely optimal for imbalanced data
120
121 # Plot F1 and F2 as a function of threshold to find the sweet
122 # spot.
123
124 best_name = max(results, key=lambda n: results[n]['ap'])
125 proba_best = results[best_name]['proba']
126 prec, rec, thresholds = precision_recall_curve(y, proba_best)
127
128 f1_scores = 2 * prec[:-1] * rec[:-1] / (prec[:-1] + rec[:-1] + 1e-9)
129 f2_scores = 5 * prec[:-1] * rec[:-1] / (4*prec[:-1] + rec[:-1] + 1e-
130
131 axes[1].plot(thresholds, f1_scores, color='#0d47a1', label='F1 score
132                  ', lw=2)
133 axes[1].plot(thresholds, f2_scores, color='#e65100', label='F2 score
134                  ', lw=2)
135 axes[1].axvline(0.5, color='gray', linestyle='--', label='Default
136                  threshold')
137
138 best_f1_thresh = thresholds[np.argmax(f1_scores)]
139 axes[1].axvline(best_f1_thresh, color='#0d47a1', linestyle=':', lw=1
140                  .5,
141                  label=f'Best F1 threshold={best_f1_thresh:.2f}')
142
143 axes[1].set_xlabel("Decision Threshold", fontsize=12)
144 axes[1].set_ylabel("Score", fontsize=12)
145 axes[1].set_title(f"F1/F2 vs Threshold {best_name}", fontsize=13,
146                  fontweight='bold')
147 axes[1].legend(fontsize=9)
148 axes[1].grid(alpha=0.3)
149 axes[1].set_xlim(0, 1)
150

```

```

142 plt.tight_layout()
143 plt.savefig("imbalanced_metrics.png", dpi=150)
144 plt.close()
145
146 # -----
147 # 6. Optimal threshold classification report
148 # -----
149 optimal_labels = (proba_best >= best_f1_thresh).astype(int)
150 print(f"\nClassification report at optimal threshold ({best_f1_
    thresh:.2f}):")
151 print(classification_report(y, optimal_labels,
    target_names=['No Fraud', 'Fraud']))
152

```

3.3.1 Reading Precision-Recall Curves

The **Precision-Recall (PR) curve** plots precision on the y-axis against recall on the x-axis as the decision threshold sweeps from 0 to 1. The **Average Precision (AP)** score is the area under this curve.

For imbalanced datasets, the PR curve is more informative than the ROC curve because it directly measures performance on the minority class. The ROC curve can appear excellent ($AUC > 0.9$) even when the model performs poorly at actually finding positive cases, because the large number of true negatives inflates the TN count.

The dashed horizontal line at the baseline class frequency is the AP score of a random classifier. Any model must substantially exceed this baseline to be useful.

3.3.2 Threshold Optimisation

The decision threshold of 0.5 is an arbitrary convention. For imbalanced problems, the optimal threshold is almost never 0.5. The threshold should be chosen based on the cost asymmetry of the problem: if false negatives are 5 times more costly than false positives, use $\beta = \sqrt{5} \approx 2.2$ in the F_β score to find the optimal threshold.

Metric	Use when	Range
Accuracy	Classes are balanced AND errors are equally costly	[0, 1]
Balanced Accuracy	Imbalanced, quick threshold-free comparison	[0, 1]
F1 Score	False positives and negatives equally costly	[0, 1]
F2 Score	Missing positives more costly than false alarms	[0, 1]
Average Precision	Need full threshold-independent evaluation	[0, 1]
ROC-AUC	Comparing discriminative ability across thresholds	[0.5, 1]

Table 3: Metric selection guide for imbalanced classification problems.

4 Model Interpretability

4.1 Why Interpretability Is Not Optional

A model that cannot be explained is a model that cannot be trusted, audited, debugged, or improved. In a business context, interpretability matters for three distinct reasons:

Regulatory compliance: Legislation such as GDPR (right to explanation), the EU AI Act, and US financial regulations often require that models provide explanations for decisions that affect individuals.

Model debugging: If your model performs worse on a demographic subgroup or in a specific time period, interpretability methods reveal which features are driving that behaviour.

Stakeholder trust: A data scientist who can say “the model flags this customer as high-risk primarily because their last payment was 60 days late and their credit utilisation is 94%” earns far more trust than one who shows a confusion matrix.

4.2 SHAP: SHapley Additive exPlanations

4.2.1 The Theory

SHAP is grounded in cooperative game theory. A **Shapley value** is the average marginal contribution of a feature across all possible orderings (coalitions) of features. For feature j in a prediction $f(\mathbf{x})$:

$$\phi_j = \sum_{S \subseteq F \setminus \{j\}} \frac{|S|! (|F| - |S| - 1)!}{|F|!} [f_{S \cup \{j\}}(\mathbf{x}_{S \cup \{j\}}) - f_S(\mathbf{x}_S)]$$

where F is the set of all features, S is any subset not containing j , and f_S is the model’s prediction using only features in S . The SHAP value ϕ_j represents the feature’s fair share of the difference between the model’s prediction and the global base rate.

The key property that makes SHAP uniquely appealing is that the SHAP values sum exactly to the model’s output relative to the expected output:

$$f(\mathbf{x}) = \mathbb{E}[f(X)] + \sum_{j=1}^{|F|} \phi_j$$

This additive decomposition means every prediction is fully accounted for, with no residual “unexplained” component.

SHAP: Global and Local Explanations

```

1 import numpy as np
2 import pandas as pd
3 import matplotlib
4 matplotlib.use('Agg')
5 import matplotlib.pyplot as plt
6 import shap

```

```

7   from sklearn.datasets import make_classification
8   from sklearn.ensemble import GradientBoostingClassifier
9   from sklearn.model_selection import train_test_split
10  from sklearn.preprocessing import StandardScaler
11  from sklearn.pipeline import Pipeline
12
13  # -----
14  # 1. Dataset with interpretable feature names
15  #     (simulating a credit scoring context)
16  # -----
17  np.random.seed(42)
18  n = 2000
19  X_raw = pd.DataFrame({
20      'days_late_last_payment' : np.random.exponential(10, n).clip(0,
21          120),
22      'credit_utilisation_pct' : np.random.beta(2, 5, n) * 100,
23      'num_open_accounts' : np.random.poisson(4, n),
24      'income_annual_k' : np.random.lognormal(4, 0.5, n),
25      'years_credit_history' : np.random.gamma(3, 3, n).clip(0.5,
26          30),
27      'num_hard_inquiries' : np.random.poisson(1.5, n),
28      'debt_to_income_ratio' : np.random.beta(2, 4, n),
29      'months_since_derog' : np.random.exponential(20, n).clip(0,
30          120),
31  })
32
33  # Logistic rule with noise creates realistic credit default target
34  logit = (
35      0.05 * X_raw['days_late_last_payment']
36      + 0.03 * X_raw['credit_utilisation_pct']
37      - 0.02 * X_raw['income_annual_k'] / 10
38      + 0.5 * X_raw['debt_to_income_ratio']
39      + 0.4 * np.random.randn(n)
40      - 3.0
41  )
42  y = (1 / (1 + np.exp(-logit)) > 0.5).astype(int)
43
44
45  # -----
46  # 2. Train a gradient boosting model
47  #     We do NOT use a Pipeline here because SHAP needs the raw
48  #     model object (not wrapped), though TreeExplainer handles this.
49  # -----
50  scaler = StandardScaler()
51  X_train_s = scaler.fit_transform(X_train)
52  X_test_s = scaler.transform(X_test)
53
54  model = GradientBoostingClassifier(
55      n_estimators=200, learning_rate=0.05,
56      max_depth=4, random_state=42
57  )
58  model.fit(X_train_s, y_train)
59
60  # -----
61  # 3. Create SHAP explainer

```

```

62 #      TreeExplainer is optimised for tree-based models ( $O(TLD^2)$ )
63 #      where
64 #      T=trees, L=leaves, D=depth) much faster than brute-force.
65 # -----
66 explainer = shap.TreeExplainer(model)
67
68 # Compute SHAP values for the test set
69 # shap_values shape: (n_test_samples, n_features) for binary
70 # For binary, index [1] gives SHAP for the positive class.
71 shap_values = explainer.shap_values(X_test_s)
72
73 # -----
74 # 4. Global explanation: Beeswarm summary plot
75 #      Each dot = one test sample. X-axis = SHAP value (impact on
76 #      model output). Colour = feature value (red=high, blue=low).
77 #      Features are sorted by mean absolute SHAP value.
78 # -----
79 shap.summary_plot(
80     shap_values, X_test,      # use un-scaled X for readable values
81     plot_type='dot',
82     show=False,
83     max_display=8
84 )
85 plt.tight_layout()
86 plt.savefig("shap_beeswarm.png", dpi=150, bbox_inches='tight')
87 plt.close()
88
89 # -----
90 # 5. Global feature importance (bar chart version)
91 # -----
92 shap.summary_plot(
93     shap_values, X_test,
94     plot_type='bar',
95     show=False,
96     max_display=8
97 )
98 plt.tight_layout()
99 plt.savefig("shap_importance.png", dpi=150, bbox_inches='tight')
100 plt.close()
101
102 # -----
103 # 6. Local explanation: a single prediction
104 #      Interpret the decision for one high-risk customer.
105 sample_idx = np.where(y_test == 1)[0][0]    # first true positive
106 sample      = X_test_s[sample_idx:sample_idx+1]
107
108 predicted_prob = model.predict_proba(sample)[0, 1]
109 print(f"\n--- Local Explanation for Sample {sample_idx} ---")
110 print(f"Predicted probability of default: {predicted_prob:.3f}")
111 print(f"Base value (mean prediction) : {explainer.expected_value:.3f}")
112 print(f"\nSHAP contribution of each feature:")
113
114 feature_names = X_raw.columns.tolist()
115 sample_shap   = shap_values[sample_idx]
116 sample_orig   = X_test.iloc[sample_idx]
117

```

```

118 contributions = pd.DataFrame({
119     'feature'      : feature_names,
120     'value'        : sample_orig.values,
121     'shap_value'   : sample_shap
122 }).sort_values('shap_value', key=abs, ascending=False)
123
124 for _, row in contributions.iterrows():
125     direction = "increases" if row['shap_value'] > 0 else "decreases"
126     print(f" {row['feature'][:30]}: {direction} risk by "
127           f"{abs(row['shap_value']):.4f} "
128           f"(feature value = {row['value']:.2f})")
129
130 total = explainer.expected_value + sample_shap.sum()
131 print(f"\nBase rate ({explainer.expected_value:.3f}) + "
132       f"sum of contributions ({sample_shap.sum():.3f}) "
133       f"= {total:.3f} (model output before sigmoid)")
134
135 # -----
136 # 7. SHAP Dependence Plot: non-linear relationships
137 #     Shows SHAP value of one feature as a function of its raw value,
138 #     coloured by an interacting feature chosen automatically by SHAP
139 #
140 shap.dependence_plot(
141     'days_late_last_payment',
142     shap_values, X_test,
143     interaction_index='auto',
144     show=False
145 )
146 plt.tight_layout()
147 plt.savefig("shap_dependence.png", dpi=150, bbox_inches='tight')
148 plt.close()
149 print("\nSHAP plots saved.")

```

4.2.2 Reading SHAP Output

The **beeswarm plot** is the primary diagnostic for global interpretability. Each row is a feature. Each dot is a sample. The position on the x-axis is the SHAP value: positive values push the prediction toward the positive class; negative values push it away. The colour represents the feature's raw value. For example: if points with high `credit_utilisation_pct` (red) have large positive SHAP values, the model has learned that high utilisation predicts default — a result that is consistent with domain knowledge, which builds confidence.

The **local waterfall** for a single prediction shows how each feature's SHAP value adds to or subtracts from the base rate to produce the final score. This is the explanation you would present to a credit officer reviewing a declined application.

4.3 LIME: Local Interpretable Model-agnostic Explanations

LIME takes a philosophically different approach. Rather than computing exact Shapley values, LIME constructs a *local linear approximation* of the model's behaviour in the neighbourhood of a specific prediction.

The algorithm works in three steps:

- Step 1. Perturb the input.** Generate N synthetic samples by randomly turning features on and off (for tabular data, this means replacing feature values with values sampled from the training distribution).
- Step 2. Weight by proximity.** Weight each synthetic sample by its distance from the original input using an exponential kernel.
- Step 3. Fit a sparse linear model.** Train a regularised linear model (typically LASSO) on the weighted synthetic samples. The coefficients of this linear model are the local explanation.

LIME Tabular Explainer

```

1 import numpy as np
2 import pandas as pd
3 import matplotlib
4 matplotlib.use('Agg')
5 import matplotlib.pyplot as plt
6 import lime
7 import lime.lime_tabular
8 from sklearn.ensemble import GradientBoostingClassifier
9 from sklearn.model_selection import train_test_split
10 from sklearn.preprocessing import StandardScaler
11
12 # Reuse the credit scoring dataset from the SHAP section
13 np.random.seed(42)
14 n = 2000
15 X_raw = pd.DataFrame({
16     'days_late_last_payment' : np.random.exponential(10, n).clip(0,
17                                                       120),
18     'credit_utilisation_pct' : np.random.beta(2, 5, n) * 100,
19     'num_open_accounts' : np.random.poisson(4, n),
20     'income_annual_k' : np.random.lognormal(4, 0.5, n),
21     'years_credit_history' : np.random.gamma(3, 3, n).clip(0.5,
22                                                       30),
23     'num_hard_inquiries' : np.random.poisson(1.5, n),
24     'debt_to_income_ratio' : np.random.beta(2, 4, n),
25     'months_since_derog' : np.random.exponential(20, n).clip(0,
26                                                       120),
27 })
28 logit = (0.05*X_raw['days_late_last_payment']
29           + 0.03*X_raw['credit_utilisation_pct']
30           - 0.02*X_raw['income_annual_k']/10
31           + 0.5*X_raw['debt_to_income_ratio']
32           + 0.4*np.random.randn(n) - 3.0)
33 y = (1/(1+np.exp(-logit)) > 0.5).astype(int)
34
35 X_train, X_test, y_train, y_test = train_test_split(
36     X_raw.values, y, test_size=0.25, random_state=42, stratify=y
37 )
38
39 scaler = StandardScaler()
40 X_tr_s = scaler.fit_transform(X_train)
41 X_te_s = scaler.transform(X_test)

```

```

40 model = GradientBoostingClassifier(
41     n_estimators=200, learning_rate=0.05, max_depth=4, random_state=
42
42 )
43 model.fit(X_tr_s, y_train)
44
45 # -----
46 # 1. Build the LIME explainer
47 #     training_data: the SCALED training set, so LIME samples from
48 #         the
49 #         correct distribution.
50 #     feature_names: column names for readable output.
51 #     class_names: label strings.
52 #     discretize_continuous: if True, LIME bins continuous features
53 #         this makes explanations more readable but less precise.
53 # -----
54 explainer = lime.lime_tabular.LimeTabularExplainer(
55     training_data=X_tr_s,
56     feature_names=X_raw.columns.tolist(),
57     class_names=['No Default', 'Default'],
58     mode='classification',
59     discretize_continuous=True,
60     random_state=42
61 )
62
63 # -----
64 # 2. Explain a single prediction
65 #     num_features: how many features to include in the local model.
66 #     num_samples: how many perturbed samples to generate (more =
67 #         stable).
67 # -----
68 idx = np.where(y_test == 1)[0][3]      # a true default case
69 sample = X_te_s[idx]
70
71 explanation = explainer.explain_instance(
72     data_row=sample,
73     predict_fn=model.predict_proba,
74     num_features=8,
75     num_samples=3000,
76     top_labels=1
77 )
78
79 pred_prob = model.predict_proba(sample.reshape(1,-1))[0,1]
80 print(f"Predicted probability of default: {pred_prob:.3f}")
81 print("\nLIME local explanation (feature: weight):")
82 for feat, weight in explanation.as_list(label=1):
83     direction = "increases risk" if weight > 0 else "decreases risk"
84     print(f"  {feat}<50: {weight:+.4f}  ({direction})")
85
86 # -----
87 # 3. Save the LIME explanation as an HTML file
88 #     This is LIME's native format and renders feature weights as
89 #         a horizontal bar chart.
90 # -----
91 html_str = explanation.as_html()
92 with open("lime_explanation.html", "w") as f:
93     f.write(html_str)
94 print("\nLIME HTML explanation saved.")

```

```

95
96  # -----
97  # 4. Manual bar chart (for embedding in reports)
98  #
99  lime_list = sorted(explanation.as_list(label=1), key=lambda x: x[1])
100 features_l = [item[0] for item in lime_list]
101 weights_l = [item[1] for item in lime_list]
102 colors_l = ['#e65100' if w > 0 else '#0d47a1' for w in weights_l]
103
104 fig, ax = plt.subplots(figsize=(9, 5))
105 bars = ax.barh(features_l, weights_l, color=colors_l, edgecolor='white')
106 ax.axvline(0, color='black', linewidth=0.8)
107 ax.set_xlabel("LIME weight (positive = increases P(Default))",
108                 fontsize=11)
109 ax.set_title(f'LIME Local Explanation Sample {idx}\n'
110               f'P(Default)={pred_prob:.3f}', fontsize=12, fontweight='bold')
111 ax.grid(axis='x', alpha=0.3)
112 plt.tight_layout()
113 plt.savefig("lime_explanation.png", dpi=150, bbox_inches='tight')
114 plt.close()
115 print("LIME bar chart saved.")

```

4.3.1 SHAP vs LIME: When to Use Each

Property	SHAP	LIME
Mathematical foundation	Shapley values (game theory)	Local linear surrogate
Consistency	Globally consistent; local explanations sum to prediction	Locally consistent only; can vary between runs
Model agnosticism	TreeExplainer requires tree models; KernelExplainer is model-agnostic but slow	Fully model-agnostic
Computational cost	Fast for trees, slow for deep models	Moderate; depends on num_samples
Stability	Deterministic (for TreeExplainer)	Stochastic; run multiple times to check variance
Best for	Audits, global analysis, regulatory documentation	Quick local explanations, non-tree models

Table 4: SHAP versus LIME comparison.

4.4 Partial Dependence Plots

Partial Dependence Plots (PDP) show the *marginal* effect of one or two features on the model's predicted outcome, averaging over all other features. Formally, for feature j :

$$\hat{f}_j(x_j) = \mathbb{E}_{X_{-j}} \left[\hat{f}(x_j, X_{-j}) \right] \approx \frac{1}{n} \sum_{i=1}^n \hat{f}(x_j, \mathbf{x}_{-j}^{(i)})$$

For each value of x_j , the model prediction is computed for every training sample (fixing x_j while using each sample's actual values for all other features), and the results are averaged. This removes the dependence on all other features, showing the isolated effect of x_j .

Partial Dependence Plots and ICE Curves

```

1  import numpy as np
2  import pandas as pd
3  import matplotlib
4  matplotlib.use('Agg')
5  import matplotlib.pyplot as plt
6  from sklearn.inspection import PartialDependenceDisplay
7  from sklearn.ensemble import GradientBoostingClassifier
8  from sklearn.model_selection import train_test_split
9  from sklearn.preprocessing import StandardScaler
10
11 np.random.seed(42)
12 n = 2000
13 X_raw = pd.DataFrame({
14     'days_late_last_payment' : np.random.exponential(10, n).clip(0,
15             120),
16     'credit_utilisation_pct' : np.random.beta(2, 5, n) * 100,
17     'num_open_accounts' : np.random.poisson(4, n),
18     'income_annual_k' : np.random.lognormal(4, 0.5, n),
19     'years_credit_history' : np.random.gamma(3, 3, n).clip(0.5,
20             30),
21     'num_hard_inquiries' : np.random.poisson(1.5, n),
22     'debt_to_income_ratio' : np.random.beta(2, 4, n),
23     'months_since_derog' : np.random.exponential(20, n).clip(0,
24             120),
25 })
26 logit = (0.05*X_raw['days_late_last_payment']
27           + 0.03*X_raw['credit_utilisation_pct']
28           - 0.02*X_raw['income_annual_k']/10
29           + 0.5*X_raw['debt_to_income_ratio']
30           + 0.4*np.random.randn(n) - 3.0)
31 y = (1/(1+np.exp(-logit)) > 0.5).astype(int)
32
33 X_train, X_test, y_train, y_test = train_test_split(
34     X_raw, y, test_size=0.25, random_state=42, stratify=y
35 )
36
37 scaler = StandardScaler()
38 X_tr_s = pd.DataFrame(scaler.fit_transform(X_train), columns=X_raw.
39     .columns)
40 X_te_s = pd.DataFrame(scaler.transform(X_test), columns=X_raw.
41     .columns)
42
43 model = GradientBoostingClassifier(
44     n_estimators=200, learning_rate=0.05, max_depth=4, random_state=
45             42
46 )
47 model.fit(X_tr_s, y_train)

```

```

42
43 # -----
44 # 1. PDP for the top 4 features
45 #     kind='both': overlay ICE (individual conditional expectation)
46 #     curves on top of the mean PDP line.
47 #     ICE curves reveal heterogeneity: if all ICE curves are parallel
48 #         ,
49 #             the feature has the same effect on everyone. If they fan out,
50 #             there are interaction effects with other features.
51 # -----
52 features_to_plot = [
53     'days_late_last_payment',
54     'credit_utilisation_pct',
55     'debt_to_income_ratio',
56     'income_annual_k'
57 ]
58 fig, axes = plt.subplots(2, 2, figsize=(13, 9))
59 axes_flat = axes.flatten()
60
61 for ax, feat in zip(axes_flat, features_to_plot):
62     disp = PartialDependenceDisplay.from_estimator(
63         model,
64         X_tr_s,
65         features=[feat],
66         kind='both',                      # PDP + ICE
67         subsample=200,                   # use 200 samples for ICE (speed)
68         n_jobs=-1,
69         grid_resolution=60,
70         ax=ax,
71         ice_lines_kw={'color': '#90caf9', 'alpha': 0.3, 'linewidth': 0.5},
72         pd_line_kw={'color': '#0d47a1', 'linewidth': 2.5, 'label': 'PDP mean'}
73     )
74     ax.set_title(f"PDP + ICE: {feat}", fontsize=11, fontweight='bold')
75     ax.set_xlabel(feat, fontsize=9)
76     ax.set_ylabel("Partial dependence", fontsize=9)
77     ax.grid(alpha=0.3)
78
79 plt.suptitle("Partial Dependence Plots with ICE Curves",
80               fontsize=14, fontweight='bold', y=1.01)
81 plt.tight_layout()
82 plt.savefig("pdp_ice.png", dpi=150, bbox_inches='tight')
83 plt.close()
84
85 # -----
86 # 2. 2D PDP: interaction between two features
87 #     If the surface is not a simple sum of two 1D functions, there
88 #     is a genuine statistical interaction between the two features.
89 # -----
90 fig, ax = plt.subplots(figsize=(8, 6))
91 disp_2d = PartialDependenceDisplay.from_estimator(
92     model,
93     X_tr_s,
94     features=[('days_late_last_payment', 'credit_utilisation_pct')],
95     kind='average',

```

```

96     grid_resolution=30,
97     ax=ax,
98     n_jobs=-1
99 )
100 ax.set_title("2D PDP: Late Payment vs Credit Utilisation",
101                 fontsize=12, fontweight='bold')
102 plt.tight_layout()
103 plt.savefig("pdp_2d.png", dpi=150, bbox_inches='tight')
104 plt.close()
105 print("PDP plots saved.")

```

Individual Conditional Expectation (ICE) curves expose something that PDPs hide. A PDP averages over all samples; if some samples see a positive effect from a feature and others see a negative effect, those effects cancel out and the PDP appears flat. ICE curves plot the predicted outcome for each individual sample as the feature varies, revealing this heterogeneity. When ICE curves cross each other, a feature interaction is present.

The **2D PDP** extends this to pairs of features. The contour plot shows the model's predicted outcome as a function of two features simultaneously. Non-linear contours that cannot be decomposed into a sum of a row pattern and a column pattern indicate an interaction effect.

Interpretability methods are not only for explainability to stakeholders — they are essential debugging tools. If a SHAP summary plot shows a feature contributing positively that you know from domain expertise should contribute negatively, you have found a bug: possibly a data leakage issue, a target encoding done incorrectly, or a labelling error. Always validate model behaviour against domain knowledge before deployment.

5 Putting It All Together: End-to-End Evaluation Workflow

This section demonstrates a complete, production-grade evaluation workflow that combines all four pillars from a single coherent script. The pattern shown here is the template you should follow for any classification problem.

Full Evaluation Pipeline — Production Template

```

1  import numpy as np
2  import pandas as pd
3  import matplotlib
4  matplotlib.use('Agg')
5  import matplotlib.pyplot as plt
6  import shap
7  import warnings
8  warnings.filterwarnings('ignore')

9
10 from sklearn.datasets import make_classification
11 from sklearn.ensemble import GradientBoostingClassifier
12 from sklearn.model_selection import (
13     StratifiedKFold, cross_validate, RandomizedSearchCV, cross_val_
14     predict
15 )
16 from sklearn.pipeline import Pipeline
17 from sklearn.preprocessing import StandardScaler
18 from sklearn.metrics import (
19     classification_report, average_precision_score,
20     roc_auc_score, balanced_accuracy_score, f1_score,
21     precision_recall_curve
22 )
23 from scipy.stats import randint, loguniform, uniform
24 print("=" * 65)
25 print(" MODEL EVALUATION & VALIDATION PRODUCTION WORKFLOW")
26 print("=" * 65)
27
28 #
29 # STEP 1: Data
30 #
31 X, y = make_classification(
32     n_samples=3000, n_features=20, n_informative=10,
33     weights=[0.88, 0.12], random_state=42
34 )
35 feature_names = [f"feature_{i:02d}" for i in range(X.shape[1])]
36 print(f"\nStep 1 Dataset: {X.shape[0]} samples, "
37       f"{X.shape[1]} features, {100*y.mean():.1f}% positive")
38
39 # Reserve a true holdout set touch this ONLY at the very end.
40 from sklearn.model_selection import train_test_split
41 X_dev, X_holdout, y_dev, y_holdout = train_test_split(
42     X, y, test_size=0.15, random_state=42, stratify=y
43 )
44 print(f"      Dev set: {X_dev.shape[0]} samples | "
45       f" Holdout: {X_holdout.shape[0]} samples")

```

```

46
47  #
48  # STEP 2: Hyperparameter tuning (inner CV loop)
49  #
50  print("\nStep 2  Hyperparameter Tuning (RandomizedSearchCV)... ")
51
52  inner_cv = StratifiedKFold(n_splits=3, shuffle=True, random_state=
53  42)
54
55  pipeline = Pipeline([
56      ('scaler', StandardScaler()),
57      ('clf', GradientBoostingClassifier(random_state=42))
58  ])
59
60  param_dist = {
61      'clf__n_estimators' : randint(50, 300),
62      'clf__max_depth'   : randint(2, 6),
63      'clf__learning_rate': loguniform(1e-2, 0.3),
64      'clf__subsample'    : uniform(0.5, 0.5),
65      'clf__min_samples_leaf': randint(5, 30),
66  }
67
68  search = RandomizedSearchCV(
69      pipeline, param_dist,
70      n_iter=30,
71      scoring='average_precision',
72      cv=inner_cv,
73      n_jobs=-1,
74      random_state=42,
75      refit=True
76  )
77  search.fit(X_dev, y_dev)
78  best_pipeline = search.best_estimator_
79
80  print(f"      Best tuning AP (inner CV): {search.best_score_:.4f}")
81  print(f"      Best params: n_est={search.best_params_['clf__n_
82  estimators']}, "
83  f"depth={search.best_params_['clf__max_depth']}, "
84  f"lr={search.best_params_['clf__learning_rate']:.4f}")
85
86  #
87  # STEP 3: Unbiased performance estimation (outer CV loop)
88  #
89  print("\nStep 3  Outer CV: Unbiased Performance Estimation...")
90
91  outer_cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=0)
92
93  # For the outer loop we re-run the search inside each fold to
94  # avoid optimism, but for speed we demonstrate with the best
95  # pipeline.
96
97  # In practice, use a nested CV (GridSearchCV inside cross_validate).
98  cv_results = cross_validate(
99      best_pipeline, X_dev, y_dev,
100     cv=outer_cv,
101     scoring={
102         'ap'       : 'average_precision',
103         'auc'     : 'roc_auc',
104         'f1'       : 'f1',
105     }
106 )

```

```

101         'bal_acc': 'balanced_accuracy',
102     },
103     return_train_score=True,
104     n_jobs=-1
105 )
106
107 print(f"\n      {'Metric':<18} {'Mean':>8} {'Std':>8} {'Overfit Gap'
108      :>12} )")
109 for m in ['ap', 'auc', 'f1', 'bal_acc']:
110     train_mean = cv_results[f'train_{m}'].mean()
111     test_mean = cv_results[f'test_{m}'].mean()
112     test_std = cv_results[f'test_{m}'].std()
113     gap = train_mean - test_mean
114     print(f"      {m:<18} {test_mean:>8.4f} {test_std:>8.4f} {gap:>
115          12.4f} )")
116 #
117 # STEP 4: Threshold optimisation
118 #
119 print("\nStep 4 Threshold Optimisation...")
120
121 oof_proba = cross_val_predict(
122     best_pipeline, X_dev, y_dev, cv=outer_cv,
123     method='predict_proba',
124 )[ :, 1]
125
126 prec, rec, thresholds = precision_recall_curve(y_dev, oof_proba)
127 # Use F2: recall matters 2x more than precision (missing positives
128 #           costly)
129 f2_scores = (5 * prec[:-1] * rec[:-1]) / (4*prec[:-1] + rec[:-1] + 1
130 e-9)
131 best_thresh_idx = np.argmax(f2_scores)
132 optimal_threshold = thresholds[best_thresh_idx]
133
134 print(f"      Optimal threshold (max F2): {optimal_threshold:.3f}")
135 print(f"      F2 at optimal: {f2_scores[best_thresh_idx]:.4f}")
136 print(f"      F2 at 0.50 : {f2_scores[np.argmin(np.abs(thresholds
137 - 0.5))]:.4f}")
138 #
139 # STEP 5: Final holdout evaluation (one-time, no peeking before this
140 #           )
141 #
142 print("\nStep 5 Final Holdout Evaluation...")
143
144 best_pipeline.fit(X_dev, y_dev)
145 holdout_proba = best_pipeline.predict_proba(X_holdout)[:, 1]
146 holdout_labels = (holdout_proba >= optimal_threshold).astype(int)
147
148 print(f"\n      *** FINAL HOLDOUT RESULTS ***")
149 print(f"      AUC-ROC : {roc_auc_score(y_holdout, holdout_
150      proba):.4f}")
151 print(f"      Average Precision: {average_precision_score(y_holdout
152      , holdout_proba):.4f}")
153 print(f"      Balanced Accuracy: {balanced_accuracy_score(y_holdout
154      , holdout_labels):.4f}")

```

```

149  print(f"      F1 Score           : {f1_score(y_holdout, holdout_
150    labels):.4f}")
150  print()
151  print(classification_report(y_holdout, holdout_labels,
152                               target_names=['Negative', 'Positive']))
153
154  #
155  # STEP 6: SHAP interpretability on holdout set
156  #
157  print("Step 6  SHAP Global Interpretability...")
158
159  fitted_clf     = best_pipeline.named_steps['clf']
160  fitted_scaler = best_pipeline.named_steps['scaler']
161  X_holdout_s   = fitted_scaler.transform(X_holdout)
162
163  explainer     = shap.TreeExplainer(fitted_clf)
164  shap_vals     = explainer.shap_values(X_holdout_s)
165
166  # Compute mean absolute SHAP and rank features
167  mean_abs_shap = np.abs(shap_vals).mean(axis=0)
168  feat_importance = pd.Series(mean_abs_shap, index=feature_names) \
169    .sort_values(ascending=False)
170
171  print("\n      Top 5 features by mean |SHAP|:")
172  for feat, val in feat_importance.head(5).items():
173      print(f"          {feat}: {val:.5f}")
174
175  shap.summary_plot(shap_vals, X_holdout,
176                    feature_names=feature_names,
177                    plot_type='bar', show=False, max_display=10)
178  plt.tight_layout()
179  plt.savefig("final_shap_importance.png", dpi=150, bbox_inches='tight')
180  plt.close()
181  print("\n      SHAP importance plot saved.")
182  print("\n" + "="*65)
183  print("  EVALUATION COMPLETE")
184  print("="*65)

```

6 The Evaluation Checklist

Before declaring a model production-ready, work through the following checklist. Each item addresses a common failure mode described in this tutorial.

Data Integrity

- All preprocessing fitted inside a Pipeline or on training data only.
- No future data in training splits (especially for time-series).
- Class distribution verified in each fold (use StratifiedKFold).
- Holdout test set segregated at the start and not inspected until the final evaluation.

Cross-Validation

- CV strategy matches deployment scenario (StratifiedKFold for classification, TimeSeriesSplit for temporal data).
- Both mean and standard deviation of CV scores reported.
- Training scores compared to validation scores to diagnose overfitting.
- Performance trend across time-series folds inspected for concept drift.

Hyperparameter Tuning

- Tuning performed on development set only; holdout never used during tuning.
- Nested cross-validation used if reporting unbiased performance alongside tuning.
- Convergence plot inspected to confirm budget was sufficient.

Metrics

- Metric chosen based on business cost structure, not convenience.
- For imbalanced datasets: accuracy not used alone; AP, F1, and balanced accuracy reported.
- Decision threshold optimised for the chosen business objective.
- Precision-recall curve inspected, not just the single-threshold result.

Interpretability

- SHAP global summary consistent with domain knowledge.
- At least one local explanation generated and reviewed for a representative positive prediction and a representative negative prediction.
- PDP or ICE curves used to verify expected monotonicity of key features.

- Any surprising model behaviour investigated for data errors or leakage before proceeding.

7 Regression Metrics: Measuring Continuous Predictions

7.1 The Problem with a Single Number

Classification has the luxury of a confusion matrix — a two-dimensional structure that separates types of errors. Regression collapses the comparison between predicted and actual values into a single error distribution. The danger is that different metrics collapse that distribution differently, each hiding a different kind of problem. A model can have a low RMSE but a terrible MAPE on small values, or a good R^2 but systematically biased predictions in one region. Understanding each metric's blind spots is what separates a careful practitioner from someone who reports the first number that looks good.

7.2 Core Regression Metrics

Let y_i be the true value and \hat{y}_i be the prediction, with n samples and mean true value \bar{y} .

Mean Absolute Error (MAE) is the average of the absolute differences. It is in the same unit as the target and gives every error equal weight:

$$\text{MAE} = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

Mean Squared Error (MSE) squares each error before averaging, which means large errors are penalised disproportionately. This is appropriate when large errors are especially costly:

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

Root Mean Squared Error (RMSE) is the square root of MSE, restoring the original unit and making it directly comparable to MAE:

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

Mean Absolute Percentage Error (MAPE) expresses error as a percentage of the true value, making it scale-independent and useful for comparing models across datasets with different magnitudes. It is undefined or explosive when $y_i \approx 0$:

$$\text{MAPE} = \frac{100}{n} \sum_{i=1}^n \left| \frac{y_i - \hat{y}_i}{y_i} \right|$$

Symmetric MAPE (sMAPE) fixes the division-by-zero problem by dividing by the mean of the actual and predicted values:

$$\text{sMAPE} = \frac{100}{n} \sum_{i=1}^n \frac{2|y_i - \hat{y}_i|}{|y_i| + |\hat{y}_i|}$$

R^2 (**Coefficient of Determination**) measures the proportion of variance in the target explained by the model. A baseline model that always predicts \bar{y} achieves $R^2 = 0$; a perfect model achieves $R^2 = 1$. Negative values mean the model is worse than the mean:

$$R^2 = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$$

Adjusted R^2 penalises for the number of features p , preventing R^2 from artificially increasing when irrelevant features are added:

$$\bar{R}^2 = 1 - (1 - R^2) \cdot \frac{n - 1}{n - p - 1}$$

Metric	Use when	Caution
MAE	Robust to outliers; median-optimising	Not differentiable at 0; less sensitive to large errors
RMSE	Large errors are disproportionately costly	Heavily influenced by outliers; unit-sensitive
MAPE	Comparing across scales; percentage makes sense	Blows up near zero; biased toward under-prediction
sMAPE	MAPE context but targets near zero exist	Can still behave oddly for very small values
R^2	Explaining variance; comparing to baseline	Scale-free but can be misleading with non-linear models
Adj. R^2	Feature selection; multivariate regression	Still a summary statistic; inspect residuals

Table 5: Regression metric selection guide.

7.3 Residual Analysis: Beyond a Single Number

No single metric captures the full picture. Residual analysis — examining the distribution and structure of errors — is the gold standard for diagnosing regression models.

Comprehensive Regression Evaluation with Residual Diagnostics

```

1 import numpy as np
2 import pandas as pd
3 import matplotlib
4 matplotlib.use('Agg')
5 import matplotlib.pyplot as plt
6 from sklearn.datasets import fetch_california_housing

```

```

7   from sklearn.ensemble import GradientBoostingRegressor
8   from sklearn.linear_model import LinearRegression
9   from sklearn.model_selection import cross_val_predict, KFold, cross_
    validate
10  from sklearn.pipeline import Pipeline
11  from sklearn.preprocessing import StandardScaler
12  from sklearn.metrics import (
13      mean_absolute_error, mean_squared_error,
14      r2_score, mean_absolute_percentage_error
15  )
16
17  # -----
18  # 1. California housing dataset: predict median house value
19  #     Target is in units of $100,000. ~20,000 samples, 8 features.
20  # -----
21  housing = fetch_california_housing(as_frame=True)
22  X, y = housing.data, housing.target
23  print(f"Dataset: {X.shape[0]} samples, {X.shape[1]} features")
24  print(f"Target range: [{y.min():.2f}, {y.max():.2f}]  "
25        f"Mean: {y.mean():.2f} Std: {y.std():.2f}")
26
27  # -----
28  # 2. Two models: linear baseline and gradient boosting
29  # -----
30  models = {
31      'Linear Regression': Pipeline([
32          ('scaler', StandardScaler()),
33          ('reg', LinearRegression())
34      ]),
35      'Gradient Boosting': Pipeline([
36          ('scaler', StandardScaler()),
37          ('reg', GradientBoostingRegressor(
38              n_estimators=200, learning_rate=0.05,
39              max_depth=5, random_state=42
40          ))
41      ])
42  }
43
44  cv = KFold(n_splits=5, shuffle=True, random_state=42)
45
46  def smape(y_true, y_pred):
47      """Symmetric Mean Absolute Percentage Error."""
48      denom = (np.abs(y_true) + np.abs(y_pred)) / 2
49      return np.mean(np.abs(y_true - y_pred) / denom) * 100
50
51  results = {}
52  for name, pipe in models.items():
53      # Out-of-fold predictions for unbiased metric computation
54      oof_pred = cross_val_predict(pipe, X, y, cv=cv, n_jobs=-1)
55
56      mae    = mean_absolute_error(y, oof_pred)
57      rmse   = np.sqrt(mean_squared_error(y, oof_pred))
58      mape   = mean_absolute_percentage_error(y, oof_pred) * 100
59      smape  = smape(y.values, oof_pred)
60      r2     = r2_score(y, oof_pred)
61
62      # Adjusted R^2
63      n, p  = X.shape

```

```

64     adj_r2 = 1 - (1 - r2) * (n - 1) / (n - p - 1)
65
66     results[name] = {
67         'oof_pred': oof_pred,
68         'MAE': mae, 'RMSE': rmse,
69         'MAPE': mape, 'sMAPE': smap,
70         'R2': r2, 'Adj_R2': adj_r2
71     }
72
73 # -----
74 # 3. Metric comparison table
75 # -----
76 print(f"\n{'Metric':<12}", end="")
77 for name in models:
78     print(f" {name:<22}", end="")
79 print()
80 print("-" * 60)
81 for metric in ['MAE', 'RMSE', 'MAPE', 'sMAPE', 'R2', 'Adj_R2']:
82     print(f"{metric:<12}", end="")
83     for name in models:
84         val = results[name][metric]
85         unit = "%" if metric in ['MAPE', 'sMAPE'] else ""
86         print(f" {val:>8.4f}{unit:<14}", end="")
87     print()
88
89 # -----
90 # 4. Four-panel residual diagnostic plot (Gold Standard)
91 #     Panel A: Residuals vs Fitted      checks heteroscedasticity
92 #     Panel B: QQ plot                  checks normality of errors
93 #     Panel C: Scale-Location          detects variance instability
94 #     Panel D: Error distribution    checks symmetry and tails
95 # -----
96 from scipy import stats
97
98 fig, axes = plt.subplots(2, 2, figsize=(13, 10))
99 best_name = 'Gradient Boosting'
100 oof = results[best_name]['oof_pred']
101 resid = y.values - oof
102 std_resid = resid / resid.std()
103
104 # Panel A: Residuals vs Fitted
105 axes[0,0].scatter(oof, resid, alpha=0.3, s=8, color='#0d47a1')
106 axes[0,0].axhline(0, color='red', linewidth=1.5, linestyle='--')
107 # Add LOWESS-style trend
108 from numpy.polynomial import polynomial as P
109 z = np.polyfit(oof, resid, 3)
110 p = np.poly1d(z)
111 x_line = np.linspace(oof.min(), oof.max(), 200)
112 axes[0,0].plot(x_line, p(x_line), color='orange', linewidth=2)
113 axes[0,0].set_xlabel("Fitted Values", fontsize=11)
114 axes[0,0].set_ylabel("Residuals", fontsize=11)
115 axes[0,0].set_title("Residuals vs Fitted\n(should be random around 0)", fontsize=11, fontweight='bold')
116 axes[0,0].grid(alpha=0.3)
117
118 # Panel B: Normal Q-Q Plot

```

```

120  (osm, osr), (slope, intercept, r) = stats.probplot(std_resid, dist='norm')
121  axes[0,1].scatter(osm, osr, alpha=0.4, s=8, color='#0d47a1')
122  axes[0,1].plot(osm, slope*np.array(osm)+intercept, 'r--', linewidth=1.5)
123  axes[0,1].set_xlabel("Theoretical Quantiles", fontsize=11)
124  axes[0,1].set_ylabel("Sample Quantiles", fontsize=11)
125  axes[0,1].set_title("Normal Q-Q Plot\n(deviations = non-normal errors)", fontsize=11, fontweight='bold')
126  axes[0,1].grid(alpha=0.3)
127
128
129 # Panel C: Scale-Location (sqrt of abs residuals vs fitted)
130 axes[1,0].scatter(oof, np.sqrt(np.abs(std_resid)), alpha=0.3, s=8, color='#00838f')
131 axes[1,0].set_xlabel("Fitted Values", fontsize=11)
132 axes[1,0].set_ylabel(r"\$\\sqrt{|\text{Std. Residuals}|}\$", fontsize=11)
133 axes[1,0].set_title("Scale-Location\n(flat line = homoscedastic)", fontsize=11, fontweight='bold')
134 axes[1,0].grid(alpha=0.3)
135
136
137 # Panel D: Error distribution
138 axes[1,1].hist(resid, bins=60, color='#0d47a1', alpha=0.75, edgecolor='white',
139                 linewidth=0.3, density=True)
140 xr = np.linspace(resid.min(), resid.max(), 300)
141 axes[1,1].plot(xr, stats.norm.pdf(xr, resid.mean(), resid.std()), 'r--', linewidth=2, label='Normal fit')
142 axes[1,1].set_xlabel("Residual", fontsize=11)
143 axes[1,1].set_ylabel("Density", fontsize=11)
144 axes[1,1].set_title(f"Error Distribution\n"
145                     f"Skew={stats.skew(resid):.3f} "
146                     f"Kurt={stats.kurtosis(resid):.3f}", fontsize=11, fontweight='bold')
147 axes[1,1].legend(fontsize=9)
148 axes[1,1].grid(alpha=0.3)
149
150
151 plt.suptitle(f"Residual Diagnostics {best_name}",
152               fontsize=14, fontweight='bold', y=1.01)
153 plt.tight_layout()
154 plt.savefig("residual_diagnostics.png", dpi=150, bbox_inches='tight')
155
156
157 plt.close()
158 print("\nResidual diagnostic plots saved.")
159
160 #
161 # 5. Prediction interval width analysis
162 #     Pinball / quantile loss: for each quantile tau, a model that
163 #     outputs the tau-quantile should have tau fraction of actuals
164 #     below.
165 #
166 from sklearn.ensemble import GradientBoostingRegressor as GBR
167
168 quantiles = [0.1, 0.5, 0.9]
169 coverage_results = {}
170
171 for q in quantiles:

```

```

171     q_model = Pipeline([
172         ('scaler', StandardScaler()),
173         ('reg', GBR(loss='quantile', alpha=q,
174                      n_estimators=150, max_depth=4, random_state=42))
175     ])
176     q_pred = cross_val_predict(q_model, X, y, cv=cv, n_jobs=-1)
177     coverage = np.mean(y.values <= q_pred)
178     coverage_results[q] = coverage
179
180     print("\nQuantile calibration (target coverage vs actual coverage):"
181          )
182     for q, cov in coverage_results.items():
183         print(f" Q{int(q*100):>2}: target={q:.2f} actual={cov:.3f} "
184               f"{'OK' if abs(q - cov) < 0.03 else 'MISCALIBRATED'}")

```

The four-panel residual plot is the equivalent of the confusion matrix for regression. The *Residuals vs Fitted* panel should show a random cloud centred at zero — any curve or funnel shape indicates the model is missing systematic structure. The *Q-Q plot* should follow the diagonal line — heavy tails mean the model is underpredicting the frequency of large errors. The *Scale-Location* plot should be flat — a rising trend means variance grows with the prediction (heteroscedasticity), which invalidates confidence intervals.

8 Model Calibration: When Probabilities Must Mean Something

8.1 The Calibration Problem

A classifier's predicted probability is not just a ranking device — in many business contexts it is used directly as a probability estimate. A credit scoring model that outputs 0.30 for a customer is implicitly claiming that 30% of customers with that profile will default. If in reality 60% of those customers default, the model is severely *miscalibrated*, and any downstream decision made using the raw probability (e.g., expected loss calculation, premium pricing) will be systematically wrong.

Calibration is the statistical property that a model's predicted probability of p should match the observed frequency of the event across all predictions in the neighbourhood of p .

Formally, a model is perfectly calibrated if:

$$P(Y = 1 \mid \hat{p} = p) = p \quad \forall p \in [0, 1]$$

8.2 The Reliability Diagram

The *reliability diagram* (also called the calibration curve) is the primary visual tool. Predictions are grouped into bins by their predicted probability; for each bin, the fraction of actual positive outcomes is plotted against the mean predicted probability. A perfectly calibrated model lies on the $y = x$ diagonal.

8.3 The Brier Score

The *Brier score* is the mean squared error of probability predictions:

$$\text{BS} = \frac{1}{n} \sum_{i=1}^n (\hat{p}_i - y_i)^2$$

It ranges from 0 (perfect) to 1 (perfectly wrong). It penalises confident wrong predictions heavily and decomposes into:

$$\text{BS} = \underbrace{\text{Reliability}}_{\text{calibration error}} - \underbrace{\text{Resolution}}_{\text{sharpness}} + \underbrace{\text{Uncertainty}}_{\text{base rate variance}}$$

A model can have a low Brier score either by being well calibrated or by being sharp (confident). Reporting both Brier score and the reliability diagram separates the two contributions.

Calibration Analysis, Platt Scaling, and Isotonic Regression

```

1 import numpy as np
2 import matplotlib
3 matplotlib.use('Agg')
```

```

4  import matplotlib.pyplot as plt
5  from sklearn.datasets import make_classification
6  from sklearn.ensemble import GradientBoostingClassifier,
    RandomForestClassifier
7  from sklearn.linear_model import LogisticRegression
8  from sklearn.calibration import CalibratedClassifierCV, calibration_
    curve
9  from sklearn.model_selection import train_test_split, cross_val_
    predict, StratifiedKFold
10 from sklearn.pipeline import Pipeline
11 from sklearn.preprocessing import StandardScaler
12 from sklearn.metrics import brier_score_loss
13
14 # -----
15 # 1. Dataset and train/test split
16 # -----
17 X, y = make_classification(
18     n_samples=5000, n_features=20, n_informative=10,
19     weights=[0.75, 0.25], random_state=42
20 )
21 X_train, X_test, y_train, y_test = train_test_split(
22     X, y, test_size=0.3, stratify=y, random_state=42
23 )
24
25 # -----
26 # 2. Gradient Boosting is often poorly calibrated because it was
27 #     optimised for ranking (log-loss), not for probability accuracy.
28 #     Random Forest is typically overconfident (predicted probs
29 #     cluster toward 0 and 1, not matching actual frequencies).
30 # -----
31 models = {
32     'Gradient Boosting (raw)': Pipeline([
33         ('scaler', StandardScaler()),
34         ('clf', GradientBoostingClassifier(
35             n_estimators=200, max_depth=4, random_state=42
36         ))
37     ]),
38     'Random Forest (raw)': Pipeline([
39         ('scaler', StandardScaler()),
40         ('clf', RandomForestClassifier(
41             n_estimators=200, random_state=42
42         ))
43     ]),
44 }
45
46 # -----
47 # 3. Calibrated versions
48 #     Platt Scaling (method='sigmoid'): fits a logistic regression
49 #     on top of the model's raw scores. Works well when the
50 #     calibration curve is S-shaped.
51 #
52 #     Isotonic Regression (method='isotonic'): fits a non-parametric
53 #     monotonic function. More flexible but needs more data to avoid
54 #     overfitting. Use when the calibration error is non-monotonic.
55 #
56 #     cv='prefit': the base model is already fitted; only the
57 #     calibration layer is fitted on the calibration set.
58 # -----

```

```

59     calibrated_models = {}
60     for name, pipe in models.items():
61         pipe.fit(X_train, y_train)
62         # Platt scaling
63         calib_platt = CalibratedClassifierCV(pipe, method='sigmoid', cv=
64             'prefit')
65         calib_platt.fit(X_test[:len(X_test)//2], y_test[:len(y_test)//2])
66         calibrated_models[name + ' + Platt'] = calib_platt
67
68         # Isotonic regression
69         calib_iso = CalibratedClassifierCV(pipe, method='isotonic', cv=
70             'prefit')
71         calib_iso.fit(X_test[:len(X_test)//2], y_test[:len(y_test)//2])
72         calibrated_models[name + ' + Isotonic'] = calib_iso
73
74     # Use the remaining half of test set for evaluation
75     X_eval, y_eval = X_test[len(X_test)//2:], y_test[len(y_test)//2:]
76
77     # -----
78     # 4. Compute calibration curves and Brier scores
79     # -----
80
81     all_models = {**models, **calibrated_models}
82     eval_results = {}
83
84     for name, model in all_models.items():
85         try:
86             proba = model.predict_proba(X_eval)[:, 1]
87         except Exception:
88             proba = model.predict_proba(X_eval)[:, 1]
89
90         frac_pos, mean_pred = calibration_curve(
91             y_eval, proba, n_bins=10, strategy='uniform'
92         )
93         brier = brier_score_loss(y_eval, proba)
94         eval_results[name] = {
95             'proba': proba, 'frac_pos': frac_pos,
96             'mean_pred': mean_pred, 'brier': brier
97         }
98
99     # -----
100    # 5. Reliability diagram: 2 x 2 grid comparing raw vs calibrated
101    # -----
102    fig, axes = plt.subplots(1, 2, figsize=(13, 6))
103
104    colors_raw = ['#0d47a1', '#00838f']
105    colors_cal = ['#1565c0', '#e65100', '#1b5e20', '#880e4f']
106
107    for ax, base_name, color_raw in zip(
108        axes, ['Gradient Boosting', 'Random Forest'], colors_raw):
109        ax.plot([0, 1], [0, 1], 'k--', linewidth=1.2, label='Perfect
110            calibration')
111
112        # Raw model
113        r = eval_results[base_name + ' (raw)']
114        ax.plot(r['mean_pred'], r['frac_pos'], 's-', color=color_raw,
115                linewidth=2, markersize=6,

```

```

113         label=f"Raw (Brier={r['brier']:.4f})")
114
115     # Platt
116     rp = eval_results[base_name + '_raw'] + Platt'
117     ax.plot(rp['mean_pred'], rp['frac_pos'], 'o-', color='#e65100',
118             linewidth=2, markersize=6,
119             label=f"Platt (Brier={rp['brier']:.4f})")
120
121     # Isotonic
122     ri = eval_results[base_name + '_raw'] + Isotonic'
123     ax.plot(ri['mean_pred'], ri['frac_pos'], '^--', color="#1b5e20",
124             linewidth=2, markersize=6,
125             label=f"Isotonic (Brier={ri['brier']:.4f})")
126
127     ax.set_xlabel("Mean Predicted Probability", fontsize=11)
128     ax.set_ylabel("Fraction of Positives", fontsize=11)
129     ax.set_title(f"Reliability Diagram {base_name.split(' ')[0]}",
130                  fontsize=12, fontweight='bold')
131     ax.legend(fontsize=8.5)
132     ax.grid(alpha=0.3)
133     ax.set_xlim(0, 1); ax.set_ylim(0, 1)
134
135 plt.tight_layout()
136 plt.savefig("calibration_curves.png", dpi=150, bbox_inches='tight')
137 plt.close()
138
139 # -----
140 # 6. Summary: Brier scores
141 #
142 print("Brier Score Summary (lower = better):")
143 for name, r in eval_results.items():
144     print(f" {name[:42]}: {r['brier']:.5f}")
145 print("\nNote: calibration only improves probability accuracy, not")
146 print("discrimination (AUC). Always report both AUC and Brier score."
147      )

```

When calibration matters most: any time probabilities feed into a downstream system that treats them as actual probabilities — expected value calculations, risk-adjusted pricing, multi-model ensembles that average probabilities, or Bayesian updating. When the model output is used only as a ranking (e.g., top-K recommendations, ranked lists), calibration is irrelevant and AUC is the right metric.

Calibration and discrimination (AUC) are orthogonal properties. You can have a perfectly discriminating model with terrible calibration (a monotonic transformation of the probabilities preserves AUC but destroys calibration). Always report both. Use AUC to assess whether the model separates classes; use the Brier score and reliability diagram to assess whether the probabilities are trustworthy.

9 Statistical Significance Testing for Model Comparison

9.1 The Core Question: Is Model A Actually Better?

When Model A achieves an AUC of 0.872 and Model B achieves 0.861 on the same cross-validation folds, the natural question is: is this difference real, or is it sampling noise? Without a statistical test, you cannot answer that question. Choosing the better-looking model on a single evaluation is like flipping a coin twice and concluding it is biased toward heads.

The challenge in model comparison is that the observations (CV fold scores) are *not independent* — they share training data. Standard two-sample t-tests assume independence and are anti-conservative (too likely to declare a significant difference) in this setting.

9.2 McNemar's Test: Comparing Classifiers on the Same Test Set

McNemar's test compares two classifiers on the same set of test examples. It asks: of the examples where the two models disagree, does one model get more right than the other? It operates on the contingency table of per-sample correct/incorrect decisions and requires no distributional assumptions.

$$\chi^2 = \frac{(n_{01} - n_{10})^2}{n_{01} + n_{10}}$$

where n_{01} is the number of examples Model A gets wrong but Model B gets right, and n_{10} is the reverse.

9.3 The Corrected Resampled t-Test

For comparing models across cross-validation folds, Nadeau and Bengio (2003) proposed a correction to the paired t-test that accounts for the overlap in training data between folds. For k -fold CV with n total samples and $n_{\text{test}} = n/k$ test samples per fold:

$$t = \frac{\bar{d}}{\sqrt{\left(\frac{1}{k} + \frac{n_{\text{test}}}{n_{\text{train}}}\right) s_d^2}}$$

where \bar{d} is the mean per-fold score difference and s_d^2 is the variance of those differences.

Statistical Model Comparison: McNemar and Corrected t-Test

```

1 import numpy as np
2 import pandas as pd
3 from scipy import stats
4 from scipy.stats import chi2
5 from sklearn.datasets import make_classification
6 from sklearn.ensemble import GradientBoostingClassifier,
```

```

    RandomForestClassifier
7   from sklearn.linear_model import LogisticRegression
8   from sklearn.model_selection import StratifiedKFold, cross_val_
     predict, cross_val_score
9   from sklearn.pipeline import Pipeline
10  from sklearn.preprocessing import StandardScaler
11  from sklearn.metrics import roc_auc_score
12  from itertools import combinations
13
14 X, y = make_classification(
15     n_samples=3000, n_features=20, n_informative=10,
16     weights=[0.7, 0.3], random_state=42
17 )
18
19 cv = StratifiedKFold(n_splits=10, shuffle=True, random_state=42)
20
21 models = {
22     'Logistic Regression': Pipeline([
23         ('sc', StandardScaler()),
24         ('clf', LogisticRegression(C=0.1, max_iter=1000))
25     ]),
26     'Random Forest': Pipeline([
27         ('sc', StandardScaler()),
28         ('clf', RandomForestClassifier(n_estimators=200, random_
             state=42))
29     ]),
30     'Gradient Boosting': Pipeline([
31         ('sc', StandardScaler()),
32         ('clf', GradientBoostingClassifier(
33             n_estimators=200, learning_rate=0.05,
34             max_depth=4, random_state=42
35         ))
36     ]),
37 }
38
39 # -----
40 # 1. Collect per-fold AUC scores and per-sample hard predictions
41 # -----
42 fold_scores = {}
43 oof_labels = {}
44 n = len(y)
45 n_splits = cv.get_n_splits()
46 n_test_per_fold = n // n_splits
47
48 for name, pipe in models.items():
49     fold_aucs = cross_val_score(pipe, X, y, cv=cv, scoring='roc_auc',
50         , n_jobs=-1)
51     oof_preds = cross_val_predict(pipe, X, y, cv=cv, n_jobs=-1) # 
             hard labels
52     fold_scores[name] = fold_aucs
53     oof_labels[name] = oof_preds
54     print(f"{name:<25}: AUC={fold_aucs.mean():.4f} +/- {fold_aucs.
             std():.4f}")
55
56 # -----
57 # 2. Corrected Resampled t-Test (Nadeau-Bengio)
58 #     Compares two models using their per-fold score differences.
59 #     The correction term accounts for the non-independence of CV

```

```

        folds.
59 # -----
60 def corrected_resampled_t_test(scores_a, scores_b, n_total, n_test):
61     """
62         Nadeau-Bengio corrected paired t-test for cross-validation
63             comparison.
64
65     Parameters
66     -----
67         scores_a, scores_b : arrays of per-fold scores
68         n_total : total number of samples in dataset
69         n_test : number of test samples per fold
70
71     Returns
72     -----
73         t_stat, p_value
74     """
75     k = len(scores_a)
76     d = scores_a - scores_b                         # per-fold differences
77     d_bar = d.mean()
78     s2_d = d.var(ddof=1)                            # sample variance
79
80     n_train = n_total - n_test
81     # Correction factor: 1/k + n_test/n_train
82     correction = (1.0 / k) + (n_test / n_train)
83     se = np.sqrt(correction * s2_d)
84
85     if se == 0:
86         return 0.0, 1.0
87     t_stat = d_bar / se
88     p_val = 2 * stats.t.sf(np.abs(t_stat), df=k - 1)
89     return t_stat, p_val
90
91 print("\n--- Corrected Resampled t-Test (Nadeau-Bengio) ---")
92 print(f'{Comparison':<45} {t-stat':>8} {p-value':>10} {'
93     Significant':>12})"
94 print("-" * 78)
95
96 model_names = list(models.keys())
97 for a, b in combinations(model_names, 2):
98     t, p = corrected_resampled_t_test(
99         fold_scores[a], fold_scores[b],
100        n_total=n, n_test=n_test_per_fold
101    )
102     sig = "YES (p<0.05)" if p < 0.05 else "no"
103     print(f"{a} vs {b}: {t:>8.3f} {p:>10.4f} {sig:>12}")
104
105 # -----
106 # 3. McNemar's Test: comparing hard-label decisions on same samples
107 #      Only meaningful for samples where the two models disagree.
108 # -----
109 def mcnemar_test(labels_a, labels_b, y_true):
110     """
111         Test whether two classifiers differ significantly in their
112             errors.
113         n01: A wrong, B right
114         n10: A right, B wrong
115     Returns chi-squared statistic and p-value (with continuity

```

```

        correction).

113 """
114 n01 = np.sum((labels_a != y_true) & (labels_b == y_true))
115 n10 = np.sum((labels_a == y_true) & (labels_b != y_true))
116 # Continuity-corrected McNemar
117 if n01 + n10 == 0:
118     return 0.0, 1.0
119 chi2_stat = (abs(n01 - n10) - 1)**2 / (n01 + n10)
120 p_val = 1 - chi2.cdf(chi2_stat, df=1)
121 return chi2_stat, p_val
122
123 print("\n--- McNemar's Test (on out-of-fold hard labels) ---")
124 print(f'{Comparison} :<45} {chi2:>8} {p-value:>10} "
125      f'{n01:>6} {n10:>6} {Significant:>12}')
126 print("-" * 95)
127
128 for a, b in combinations(model_names, 2):
129     chi2_val, p = mcnemar_test(oof_labels[a], oof_labels[b], y)
130     n01 = np.sum((oof_labels[a] != y) & (oof_labels[b] == y))
131     n10 = np.sum((oof_labels[a] == y) & (oof_labels[b] != y))
132     sig = "YES (p<0.05)" if p < 0.05 else "no"
133     print(f'{a} vs {b} :<20} {chi2_val:>8.3f} {p:>10.4f} "
134           f'{n01:>6} {n10:>6} {sig:>12}')
135
136 # -----
137 # 4. Multiple comparisons correction (Bonferroni)
138 # When testing k pairs, the chance of a false positive rises.
139 # Bonferroni adjusts the significance threshold to alpha/k.
140 #
141 n_comparisons = len(list(combinations(model_names, 2)))
142 alpha_bonf = 0.05 / n_comparisons
143 print(f"\nBonferroni-corrected significance threshold for {n_
144   comparisons} "
145       f"comparisons: {alpha_bonf:.4f}")

```

If you test 20 model variants and declare significance at $p < 0.05$, you expect roughly one false positive by chance alone. Always apply a correction: Bonferroni (divide threshold by number of tests, conservative) or Benjamini-Hochberg (controls the false discovery rate, less conservative). Report corrected p-values whenever comparing more than two models.

10 Bias & Fairness Evaluation

10.1 Why Fairness Is a Measurement Problem

Fairness is not a single property — it is a family of mathematical definitions that are often mutually contradictory. The practitioner’s job is to: (1) understand which definition of fairness is appropriate for the business and regulatory context, (2) measure it, and (3) make explicit trade-offs when definitions conflict.

Before any fairness analysis, a critical distinction must be made. A **protected attribute** is a feature (race, gender, age, disability) that must not be the basis for a decision. Such features may be excluded from the model, but they must still be used to *evaluate* whether the model produces disparate outcomes.

10.2 Core Fairness Definitions

Demographic Parity (also: statistical parity): the model’s positive prediction rate should be equal across groups. Used in hiring and lending where equal rates of opportunity are required.

$$P(\hat{Y} = 1 \mid A = 0) = P(\hat{Y} = 1 \mid A = 1)$$

Equalized Odds: both the true positive rate (recall) and false positive rate must be equal across groups. Used in criminal justice and medical screening where errors of both types must not be concentrated in one group.

$$P(\hat{Y} = 1 \mid Y = y, A = 0) = P(\hat{Y} = 1 \mid Y = y, A = 1) \quad \forall y \in \{0, 1\}$$

Equal Opportunity: a relaxation of equalized odds that only requires equal recall (true positive rate). Appropriate when false negatives are the primary concern.

Disparate Impact Ratio: a ratio metric used in US employment law (the “80% rule”). If the positive rate for the disadvantaged group is less than 80% of the advantaged group’s rate, disparate impact is indicated.

$$\text{DIR} = \frac{P(\hat{Y} = 1 \mid A = 1)}{P(\hat{Y} = 1 \mid A = 0)} \quad (\text{should be } \geq 0.8)$$

Fairness Audit: Measuring and Visualising Bias

```

1 import numpy as np
2 import pandas as pd
3 import matplotlib
4 matplotlib.use('Agg')
5 import matplotlib.pyplot as plt
6 from sklearn.datasets import make_classification
7 from sklearn.ensemble import GradientBoostingClassifier
8 from sklearn.model_selection import train_test_split,
9     StratifiedKFold, cross_val_predict
9 from sklearn.pipeline import Pipeline

```

```

10  from sklearn.preprocessing import StandardScaler
11  from sklearn.metrics import (confusion_matrix, roc_auc_score,
12                                precision_score, recall_score)
13
14  np.random.seed(42)
15  n = 4000
16
17  # -----
18  # 1. Simulated loan application dataset
19  #     Group A=0: majority group (60%)
20  #     Group A=1: minority group (40%)
21  #     Introduce historical bias: minority group has artificially
22  #     lower income in training data (reflecting past discrimination).
23  #
24  group = np.random.binomial(1, 0.4, n)    # 0=majority, 1=minority
25
26  income    = np.where(group == 0,
27                        np.random.normal(65, 15, n),      # majority: mean
28                        np.random.normal(52, 14, n))    # minority: mean
29
30  credit    = np.random.normal(680, 80, n) - 30 * group
31  dti       = np.random.beta(2, 5, n) + 0.05 * group
32  history   = np.random.gamma(4, 2, n) - 1.5 * group
33
34  X = np.column_stack([income, credit, dti, history,
35                        np.random.randn(n, 6)])    # 6 noise features
36  feature_names = ['income', 'credit_score', 'dti', 'credit_history',
37                     *[f'noise_{i}' for i in range(6)]]
38
39  # True default driven by financial fundamentals (fair ground truth)
40  logit = -0.04*income + 0.005*(750-credit) + 2*dti - 0.05*history
41  y      = (1/(1+np.exp(-logit + np.random.randn(n)*0.5)) > 0.5).astype
42  (int)
43
44  print(f"Default rate majority: {y[group==0].mean():.3f} "
45        f"minority: {y[group==1].mean():.3f}")
46
47  # -----
48  # 2. Train a model WITHOUT group as a feature (group-blind)
49  #     Fairness analysis uses group ONLY in evaluation, not training.
50  #
51  X_train, X_test, y_train, y_test, g_train, g_test = train_test_split
52  (
53      X, y, group, test_size=0.3, stratify=y, random_state=42
54  )
55
56  pipeline = Pipeline([
57      ('scaler', StandardScaler()),
58      ('clf', GradientBoostingClassifier(
59          n_estimators=200, max_depth=4,
60          learning_rate=0.05, random_state=42
61      ))
62  ])
63  pipeline.fit(X_train, y_train)
64  proba_test = pipeline.predict_proba(X_test)[:, 1]
65  pred_test = pipeline.predict(X_test)

```

```

64
65  # -----
66  # 3. Per-group metrics
67  # -----
68 def group_metrics(y_true, y_pred, y_proba, groups):
69     """Compute fairness-relevant metrics per group."""
70     rows = []
71     for g_val in sorted(np.unique(groups)):
72         mask = groups == g_val
73         yt = y_true[mask]
74         yp = y_pred[mask]
75         ypr = y_proba[mask]
76         cm = confusion_matrix(yt, yp, labels=[0,1])
77         tn, fp, fn, tp = cm.ravel()
78         rows.append({
79             'group' : g_val,
80             'n' : mask.sum(),
81             'pos_rate' : yp.mean(),           # demographic parity
82                         metric
83             'tpr' : tp/(tp+fn+1e-9),      # recall / equal
84                         opportunity
85             'fpr' : fp/(fp+tn+1e-9),      # false positive
86                         rate
87             'precision' : precision_score(yt, yp, zero_division=0),
88             'auc' : roc_auc_score(yt, ypr),
89             'false_neg_rate' : fn/(fn+tp+1e-9),
90         })
91     return pd.DataFrame(rows)
92
93
94 gm = group_metrics(y_test, pred_test, proba_test, g_test)
95 print("\n--- Per-Group Metrics ---")
96 print(gm.set_index('group').round(4).to_string())
97
98 # -----
99 # 4. Fairness metrics derived from group metrics
100 # -----
101 g0 = gm[gm['group']==0].iloc[0]
102 g1 = gm[gm['group']==1].iloc[0]
103
104 print("\n--- Fairness Metrics ---")
105 dir_ratio = g1['pos_rate'] / (g0['pos_rate'] + 1e-9)
106 demo_gap = abs(g0['pos_rate'] - g1['pos_rate'])
107 tpr_gap = abs(g0['tpr'] - g1['tpr'])
108 fpr_gap = abs(g0['fpr'] - g1['fpr'])
109 eq_odds = max(tpr_gap, fpr_gap)
110
111 print(f"  Disparate Impact Ratio (>=0.8 required) : {dir_ratio:.4f}")
112
113         f"{'PASS' if dir_ratio >= 0.8 else 'FAIL'}"
114 print(f"  Demographic Parity Gap (|g0 - g1|) : {demo_gap:.4f} "
115         f"{'OK' if demo_gap < 0.05 else 'CONCERN'}")
116 print(f"  Equal Opportunity Gap (TPR diff) : {tpr_gap:.4f} "
117         f"{'OK' if tpr_gap < 0.05 else 'CONCERN'}")
118 print(f"  Equalized Odds Violation (max gap) : {eq_odds:.4f} "
119         f"{'OK' if eq_odds < 0.05 else 'CONCERN'}")
120
121 # -----

```

```

117 # 5. Threshold adjustment for equalized odds
118 #     Use a different decision threshold per group to equalise TPR.
119 #     This is the most common post-processing fairness intervention.
120 #
121 from sklearn.metrics import roc_curve
122
123 thresholds_explored = np.linspace(0.1, 0.9, 200)
124
125 # Find threshold per group that achieves TPR close to target
126 target_tpr = 0.70    # desired recall for positive class in both
                      # groups
127
128 best_thresh = {}
129 for g_val in [0, 1]:
130     mask = g_test == g_val
131     yt    = y_test[mask]
132     ypr   = proba_test[mask]
133     tprs = [recall_score(yt, (ypr >= t).astype(int), zero_division=
134                 0)
135                 for t in thresholds_explored]
136     best_idx = np.argmin(np.abs(np.array(tprs) - target_tpr))
137     best_thresh[g_val] = thresholds_explored[best_idx]
138
139 print(f"\n  Threshold adjustment for TPR={target_tpr}:")
140 print(f"    Group 0 threshold: {best_thresh[0]:.3f}")
141 print(f"    Group 1 threshold: {best_thresh[1]:.3f}")
142
143 adj_pred = np.where(
144     g_test == 0,
145     (proba_test >= best_thresh[0]).astype(int),
146     (proba_test >= best_thresh[1]).astype(int)
147 )
148
149 gm_adj = group_metrics(y_test, adj_pred, proba_test, g_test)
150 print("\n--- After Threshold Adjustment ---")
151 print(gm_adj.set_index('group')[['pos_rate', 'tpr', 'fpr', 'precision']]
152       .round(4))
153
154 #
155 fig, axes = plt.subplots(1, 2, figsize=(12, 5))
156
157 metrics_to_plot = ['pos_rate', 'tpr', 'fpr', 'precision', 'auc']
158 labels = ['Pos. Rate', 'TPR', 'FPR', 'Precision', 'AUC']
159
160 for ax, (gm_df, title) in zip(axes, [
161     (gm,      'Before Adjustment'),
162     (gm_adj, 'After Threshold Adjustment')
163 ]):
164     x = np.arange(len(metrics_to_plot))
165     w = 0.35
166     vals0 = [gm_df[gm_df['group']==0].iloc[0][m] for m in metrics_to_
167             _plot]
168     vals1 = [gm_df[gm_df['group']==1].iloc[0][m] for m in metrics_to_
169             _plot]
170
171     ax.bar(x - w/2, vals0, w, label='Group 0 (majority)',


```

```

170         color='#0d47a1', alpha=0.85)
171     ax.bar(x + w/2, vals1, w, label='Group 1 (minority)',
172             color='#e65100', alpha=0.85)
173     ax.set_xticks(x); ax.set_xticklabels(labels, fontsize=10)
174     ax.set_ylim(0, 1.1)
175     ax.axhline(0.8, color='red', linestyle='--', linewidth=1,
176                 label='80% rule threshold')
177     ax.set_title(title, fontsize=12, fontweight='bold')
178     ax.legend(fontsize=8); ax.grid(axis='y', alpha=0.3)
179
180 plt.suptitle("Fairness Metrics Before and After Threshold Adjustment",
181               fontsize=13, fontweight='bold')
182 plt.tight_layout()
183 plt.savefig("fairness_metrics.png", dpi=150, bbox_inches='tight')
184 plt.close()
185 print("\nFairness visualisation saved.")

```

Chouldechova (2017) and Kleinberg et al. (2017) independently proved that demographic parity, equalized odds, and predictive parity (precision equality) *cannot all be satisfied simultaneously* unless base rates are equal across groups. Fairness analysis requires an explicit choice of which criterion to prioritise, determined by the ethical and legal context of the application, not by the algorithm.

11 Production Monitoring & Drift Detection

11.1 The Life Cycle Does Not End at Deployment

A model that performs well on the test set at deployment time may degrade over weeks or months as the world changes. This degradation has two root causes.

Data drift (or covariate shift) occurs when the distribution of input features $P(X)$ changes while the relationship $P(Y | X)$ remains stable. Example: a pandemic shifts the income distribution of loan applicants, but the relationship between income and default probability is unchanged.

Concept drift occurs when $P(Y | X)$ itself changes — the fundamental relationship between features and target has shifted. Example: a new regulation changes what constitutes fraud, so past labelling conventions no longer apply.

Both require monitoring, but they call for different responses: data drift may be handled by retraining on new data; concept drift may require redesigning the feature set or the label definition.

11.2 Population Stability Index (PSI)

The **Population Stability Index** is the most widely used metric for monitoring input drift in production. It was developed by credit risk practitioners and measures the divergence between a reference distribution and a current distribution.

$$\text{PSI} = \sum_{b=1}^B (q_b - p_b) \ln \left(\frac{q_b}{p_b} \right)$$

where p_b is the proportion of reference data in bin b and q_b is the corresponding proportion in the current data. PSI is a symmetrised version of the KL divergence. The standard thresholds are: $\text{PSI} < 0.1$ (stable), $0.1 \leq \text{PSI} < 0.2$ (minor shift, monitor closely), $\text{PSI} \geq 0.2$ (significant shift, investigate and consider retraining).

PSI, KS Drift Detection, and Performance Monitoring

```

1 import numpy as np
2 import pandas as pd
3 import matplotlib
4 matplotlib.use('Agg')
5 import matplotlib.pyplot as plt
6 from scipy import stats
7 from sklearn.datasets import make_classification
8 from sklearn.ensemble import GradientBoostingClassifier
9 from sklearn.model_selection import train_test_split
10 from sklearn.pipeline import Pipeline
11 from sklearn.preprocessing import StandardScaler
12 from sklearn.metrics import roc_auc_score
13
14 # -----
15 # 1. Simulate production lifecycle
16 #     Month 0-2: stable data (training distribution)

```

```

17 #      Month 3-5: gradual covariate shift (data drift)
18 #      Month 6-8: concept drift (relationship changes)
19 #
20 np.random.seed(42)
21
22 def generate_batch(n, drift_factor=0.0, concept_drift=False):
23     """Generate a batch of data with optional drift."""
24     X = np.random.randn(n, 10)
25     X[:, 0] += drift_factor * 2          # feature 0 shifts with drift
26     X[:, 1] += drift_factor * 1.5        # feature 1 also shifts
27
28     if not concept_drift:
29         # Original relationship: feature 0 and 2 predict y
30         logit = 1.5*X[:,0] + X[:,2] + 0.3*np.random.randn(n)
31     else:
32         # Concept drift: feature 0 now has OPPOSITE sign
33         logit = -1.5*X[:,0] + X[:,2] + 0.3*np.random.randn(n)
34
35     y = (1/(1 + np.exp(-logit)) > 0.5).astype(int)
36     return X, y
37
38 # Training data (reference)
39 X_train, y_train = generate_batch(3000, drift_factor=0.0)
40
41 # Production batches (one per month)
42 months = list(range(9))
43 monthly_data = {
44     m: generate_batch(
45         300,
46         drift_factor=max(0, (m-2)*0.4),           # drift starts month
47         3
48         concept_drift=(m >= 6)                  # concept drift
49         from month 6
50     )
51     for m in months
52 }
53
54 # -----
55 # 2. Train model on reference data
56 #
57 pipeline = Pipeline([
58     ('scaler', StandardScaler()),
59     ('clf', GradientBoostingClassifier(n_estimators=150, random_
60                                         state=42))
61 ])
62 pipeline.fit(X_train, y_train)
63
64 # -----
65 # 3. PSI function
66 #
67 def compute_psi(reference, current, n_bins=10):
68     """
69         Compute Population Stability Index between two samples.
70         Uses percentile-based bins from the reference distribution.
71     """
72     breakpoints = np.percentile(reference, np.linspace(0, 100, n_
73                                         bins+1))
74     breakpoints = np.unique.breakpoints) # remove duplicates

```

```

71
72     ref_counts = np.histogram(reference, bins=breakpoints)[0]
73     cur_counts = np.histogram(current,    bins=breakpoints)[0]
74
75     # Convert to proportions, add small epsilon to avoid log(0)
76     eps = 1e-6
77     ref_pct = ref_counts / len(reference) + eps
78     cur_pct = cur_counts / len(current)    + eps
79
80     psi = np.sum((cur_pct - ref_pct) * np.log(cur_pct / ref_pct))
81     return psi
82
83 # -----
84 # 4. Month-by-month monitoring
85 # -----
86 reference_feature = X_train[:, 0]      # monitor feature 0
87 monitoring = []
88
89 for m in months:
90     X_m, y_m = monthly_data[m]
91     proba_m = pipeline.predict_proba(X_m)[:, 1]
92
93     # PSI on feature 0
94     psi_f0 = compute_psi(reference_feature, X_m[:, 0])
95
96     # KS test on feature 0 (alternative to PSI)
97     ks_stat, ks_pval = stats.ks_2samp(reference_feature, X_m[:, 0])
98
99     # AUC requires ground truth labels (only available with delay)
100    auc = roc_auc_score(y_m, proba_m)
101
102    # Score distribution PSI (model output drift)
103    ref_scores = pipeline.predict_proba(X_train[:300])[:, 1]
104    psi_score = compute_psi(ref_scores, proba_m)
105
106    monitoring.append({
107        'month': m,
108        'psi_feature_0': psi_f0,
109        'ks_statistic': ks_stat,
110        'ks_pvalue':    ks_pval,
111        'psi_score':   psi_score,
112        'auc':         auc,
113    })
114
115 mon_df = pd.DataFrame(monitoring).set_index('month')
116
117 print("Monthly Monitoring Dashboard:")
118 print(mon_df.round(4).to_string())
119
120 # -----
121 # 5. Monitoring dashboard plot
122 # -----
123 fig, axes = plt.subplots(2, 2, figsize=(13, 9))
124
125 # PSI feature
126 ax = axes[0,0]
127 colors = ['#e65100' if v >= 0.2 else '#f57c00' if v >= 0.1
128           else '#1b5e20' for v in mon_df['psi_feature_0']]

```

```

129  ax.bar(mon_df.index, mon_df['psi_feature_0'], color=colors, alpha=0.
          85)
130  ax.axhline(0.1, color='orange', linestyle='--', linewidth=1.5, label
              ='PSI=0.1 (caution)')
131  ax.axhline(0.2, color='red',      linestyle='--', linewidth=1.5, label
              ='PSI=0.2 (alert)')
132  ax.set_title("PSI Feature 0 (Data Drift)", fontweight='bold')
133  ax.set_xlabel("Month"); ax.set_ylabel("PSI")
134  ax.legend(fontsize=8); ax.grid(alpha=0.3)
135  ax.set_xticks(months)
136
137  # KS statistic
138  ax = axes[0,1]
139  ax.plot(mon_df.index, mon_df['ks_statistic'],
           marker='o', color='#0d47a1', linewidth=2)
140  ax.axhline(0.1, color='red', linestyle='--', label='KS=0.1 (alert)')
141  ax.fill_between(mon_df.index, 0, mon_df['ks_statistic'],
                    alpha=0.15, color='#0d47a1')
142  ax.set_title("KS Statistic Feature 0", fontweight='bold')
143  ax.set_xlabel("Month"); ax.set_ylabel("KS Statistic")
144  ax.legend(fontsize=8); ax.grid(alpha=0.3)
145  ax.set_xticks(months)
146
147  # Score distribution PSI
148  ax = axes[1,0]
149  ax.bar(mon_df.index, mon_df['psi_score'], color='#00838f', alpha=0.
          85)
150  ax.axhline(0.1, color='orange', linestyle='--', linewidth=1.5)
151  ax.axhline(0.2, color='red',      linestyle='--', linewidth=1.5)
152  ax.set_title("PSI Model Score Distribution", fontweight='bold')
153  ax.set_xlabel("Month"); ax.set_ylabel("PSI")
154  ax.grid(alpha=0.3); ax.set_xticks(months)
155
156  # AUC over time
157  ax = axes[1,1]
158  ax.plot(mon_df.index, mon_df['auc'],
           marker='s', color='#1b5e20', linewidth=2.5, markersize=8)
159  ax.fill_between(mon_df.index, mon_df['auc'].iloc[0] - 0.03,
                    mon_df['auc'], alpha=0.2, color='#1b5e20')
160  ax.axhline(mon_df['auc'].iloc[0] * 0.95, color='red', linestyle='--',
               ,
               label='5% degradation threshold')
161  ax.set_title("Model AUC Over Time", fontweight='bold')
162  ax.set_xlabel("Month"); ax.set_ylabel("AUC")
163  ax.legend(fontsize=9); ax.grid(alpha=0.3)
164  ax.set_xticks(months)
165
166  plt.suptitle("Production Monitoring Dashboard", fontsize=14,
                 fontweight='bold', y=1.01)
167  plt.tight_layout()
168  plt.savefig("monitoring_dashboard.png", dpi=150, bbox_inches='tight',
               )
169  plt.close()
170
171  print("\nMonitoring dashboard saved.")

```

The monitoring dashboard separates the *early warning signals* (PSI on raw features, visible immediately in production) from the *lagging performance signal* (AUC, which requires

ground truth labels that may arrive weeks later). The PSI on the model's output score distribution is the most valuable early warning: it detects shifts in how the model is scoring the population even before labels arrive.

In a real deployment, true labels are almost never available immediately. Monitor the model's score distribution using PSI on a weekly or daily cadence. When $\text{PSI} \geq 0.2$, trigger a manual review. When $\text{PSI} \geq 0.2$ and labels arrive confirming AUC degradation, trigger retraining. Build this alerting into your MLOps pipeline from day one.

12 Learning Curves & Diagnosing Underfitting vs Overfitting

12.1 The Bias-Variance Trade-off in Practice

Every model lives on a spectrum between two failure modes.

High bias (underfitting): the model is too simple to capture the patterns in the data. Both training error and test error are high. The model makes systematic errors — the same type of mistake regardless of how much data it sees. Solution: increase model complexity or engineer better features.

High variance (overfitting): the model is too complex and has memorised the training data. Training error is low but test error is high. Performance improves with more data. Solution: regularisation, dropout, pruning, or more data.

Learning curves are the primary tool for diagnosing which failure mode dominates.

12.2 Reading Learning Curves

A **learning curve** plots training and validation performance as a function of training set size. The diagnosis is visual:

- **Both curves converge to a high error plateau:** high bias. More data will not help. The model needs to be more complex.
- **Large gap between training and validation curves:** high variance. The model would benefit from more data. If data cannot be obtained, apply regularisation.
- **Curves converge to a low error:** good fit. The model has learned the signal without overfitting.

Learning Curves and Complexity Curves

```

1 import numpy as np
2 import matplotlib
3 matplotlib.use('Agg')
4 import matplotlib.pyplot as plt
5 from sklearn.datasets import make_classification
6 from sklearn.ensemble import GradientBoostingClassifier
7 from sklearn.linear_model import LogisticRegression
8 from sklearn.tree import DecisionTreeClassifier
9 from sklearn.model_selection import learning_curve, validation_curve
10, StratifiedKFold
11 from sklearn.pipeline import Pipeline
12 from sklearn.preprocessing import StandardScaler
13 X, y = make_classification(
14     n_samples=3000, n_features=20, n_informative=8,
15     n_redundant=4, weights=[0.65, 0.35], random_state=42
16 )
17 cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

```

```

19
20 # -----
21 # 1. Learning curves for three models
22 #     train_sizes: fractions of training data to use in each step.
23 #     Returns arrays of shape (n_steps, n_cv_folds).
24 # -----
25 models = {
26     'Logistic Regression\n(High Bias)': Pipeline([
27         ('sc', StandardScaler()),
28         ('clf', LogisticRegression(C=0.001, max_iter=1000)) # over-
29             regularised
30     ]),
31     'Gradient Boosting\n(Good Fit)': Pipeline([
32         ('sc', StandardScaler()),
33         ('clf', GradientBoostingClassifier(
34             n_estimators=100, max_depth=3,
35             learning_rate=0.1, random_state=42
36         ))
37     ]),
38     'Decision Tree\n(High Variance)': Pipeline([
39         ('sc', StandardScaler()),
40         ('clf', DecisionTreeClassifier(max_depth=None, random_state=
41             42))
42     ]),
43     train_sizes = np.linspace(0.05, 1.0, 15)
44
45 fig, axes = plt.subplots(1, 3, figsize=(15, 5), sharey=True)
46
47 for ax, (name, pipe) in zip(axes, models.items()):
48     sizes, train_scores, val_scores = learning_curve(
49         pipe, X, y,
50         train_sizes=train_sizes,
51         cv=cv,
52         scoring='roc_auc',
53         n_jobs=-1,
54         return_times=False
55     )
56
57     # Mean and std across folds
58     train_mean = train_scores.mean(axis=1)
59     train_std = train_scores.std(axis=1)
60     val_mean = val_scores.mean(axis=1)
61     val_std = val_scores.std(axis=1)
62
63     ax.plot(sizes, train_mean, 'o-', color='#0d47a1', linewidth=2,
64             markersize=5, label='Training AUC')
65     ax.fill_between(sizes, train_mean - train_std,
66                     train_mean + train_std, alpha=0.15, color='#0
67                     d47a1')
68
68     ax.plot(sizes, val_mean, 's-', color='#e65100', linewidth=2,
69             markersize=5, label='Validation AUC')
70     ax.fill_between(sizes, val_mean - val_std,
71                     val_mean + val_std, alpha=0.15, color='#e65100')
72
73     gap = (train_mean[-1] - val_mean[-1])

```

```

74     ax.set_title(f"\nFinal gap: {gap:.3f}", fontsize=11,
75                 fontweight='bold')
76     ax.set_xlabel("Training set size", fontsize=10)
77     ax.legend(fontsize=9)
78     ax.grid(alpha=0.3)
79     ax.set_ylim(0.4, 1.05)
80
81 axes[0].set_ylabel("AUC", fontsize=11)
82 plt.suptitle("Learning Curves: High Bias vs Good Fit vs High
83               Variance",
84             fontsize=13, fontweight='bold', y=1.02)
85 plt.tight_layout()
86 plt.savefig("learning_curves.png", dpi=150, bbox_inches='tight')
87 plt.close()
88
89 # -----
90 # 2. Validation curve (Complexity Curve)
91 #     Vary a single hyperparameter and observe how training and
92 #     validation performance change. This reveals the sweet spot
93 #     between underfitting and overfitting for that parameter.
94 # -----
95 param_range = np.arange(1, 10)    # max_depth from 1 to 9
96
97 pipe_dt = Pipeline([
98     ('sc', StandardScaler()),
99     ('clf', DecisionTreeClassifier(random_state=42))
100 ])
101
102 train_scores_vc, val_scores_vc = validation_curve(
103     pipe_dt, X, y,
104     param_name='clf__max_depth',
105     param_range=param_range,
106     cv=cv,
107     scoring='roc_auc',
108     n_jobs=-1
109 )
110
111 fig, ax = plt.subplots(figsize=(9, 5))
112
113 tr_m = train_scores_vc.mean(axis=1)
114 tr_s = train_scores_vc.std(axis=1)
115 va_m = val_scores_vc.mean(axis=1)
116 va_s = val_scores_vc.std(axis=1)
117
118 ax.plot(param_range, tr_m, 'o-', color='#0d47a1', lw=2, label='
119   Training AUC')
120 ax.fill_between(param_range, tr_m - tr_s, tr_m + tr_s,
121                 alpha=0.15, color='#0d47a1')
122 ax.plot(param_range, va_m, 's-', color='#e65100', lw=2, label='
123   Validation AUC')
124 ax.fill_between(param_range, va_m - va_s, va_m + va_s,
125                 alpha=0.15, color='#e65100')
126
127 best_depth = param_range[np.argmax(va_m)]
128 ax.axvline(best_depth, color='green', linestyle=':', lw=1.5,
129             label=f'Best depth={best_depth}')
130
131 ax.set_xlabel("max_depth", fontsize=12)

```

```

129  ax.set_ylabel("AUC", fontsize=12)
130  ax.set_title("Validation Curve Decision Tree max_depth\n"
131          "(Left=underfit, Right=overfit, Peak=sweet spot)",
132          fontsize=12, fontweight='bold')
133  ax.legend(fontsize=10)
134  ax.grid(alpha=0.3)
135  ax.set_xticks(param_range)
136  plt.tight_layout()
137  plt.savefig("validation_curve.png", dpi=150, bbox_inches='tight')
138  plt.close()
139  print("Learning and validation curves saved.")
140
141 # -----
142 # 3. Diagnosis summary
143 # -----
144 print("\n--- Bias-Variance Diagnosis ---")
145 for name, pipe in models.items():
146     sizes, tr, val = learning_curve(
147         pipe, X, y, train_sizes=[0.8, 1.0],
148         cv=cv, scoring='roc_auc', n_jobs=-1
149     )
150     tr_final = tr[-1].mean()
151     val_final = val[-1].mean()
152     gap = tr_final - val_final
153     if tr_final < 0.75 and val_final < 0.75:
154         diagnosis = "HIGH BIAS (underfit)"
155     elif gap > 0.10:
156         diagnosis = "HIGH VARIANCE (overfit)"
157     else:
158         diagnosis = "Good fit"
159     clean = name.replace('\n', ' ')
160     print(f" {clean[:35]}: train={tr_final:.3f} val={val_final:.3f}\n"
161           f"gap={gap:.3f} -> {diagnosis}")

```

12.3 The Validation Curve: Finding the Complexity Sweet Spot

The ***validation curve*** is the natural complement to the learning curve. Instead of varying the amount of data, it varies a single hyperparameter controlling complexity. The training AUC always increases with complexity; the validation AUC rises, peaks, and then falls as the model begins to overfit. The peak of the validation curve is the optimal complexity setting — this is what hyperparameter tuning is searching for, but the validation curve makes it visually explicit.

13 Multiclass & Multi-label Metrics

13.1 Beyond Binary Classification

Many real problems involve more than two classes: document categorisation, product defect type classification, disease subtype identification, sentiment with multiple levels. The metrics from binary classification extend to multiclass settings, but the averaging strategy becomes a critical decision that changes what is being measured.

13.2 Averaging Strategies

Given per-class precision P_c , recall R_c , and F_1^c for each class $c \in \{1, \dots, K\}$:

Macro averaging: unweighted mean across classes. Every class contributes equally regardless of size. Use this when you care equally about rare and common classes.

$$F_1^{\text{macro}} = \frac{1}{K} \sum_{c=1}^K F_1^c$$

Weighted averaging: weighted mean by class support (number of true instances). Reflects the contribution of each class proportionally. Use this when larger classes matter more.

$$F_1^{\text{weighted}} = \frac{1}{\sum_c n_c} \sum_{c=1}^K n_c \cdot F_1^c$$

Micro averaging: aggregates TP, FP, FN across all classes before computing. For precision and recall, micro-average equals accuracy when all predictions are considered.

13.3 Cohen's Kappa: Agreement Beyond Chance

Cohen's Kappa measures agreement between predicted and actual labels, corrected for what would be expected by chance. It ranges from -1 (perfect disagreement) through 0 (chance agreement) to 1 (perfect agreement):

$$\kappa = \frac{p_o - p_e}{1 - p_e}$$

where p_o is the observed accuracy and p_e is the expected accuracy under chance (computed from marginal distributions). Kappa is especially useful for imbalanced multiclass problems where accuracy is misleading.

Comprehensive Multiclass Evaluation

```

1 import numpy as np
2 import pandas as pd
3 import matplotlib
4 matplotlib.use('Agg')
5 import matplotlib.pyplot as plt

```

```

6  import seaborn as sns
7  from sklearn.datasets import make_classification
8  from sklearn.ensemble import GradientBoostingClassifier
9  from sklearn.model_selection import StratifiedKFold, cross_val_
   predict
10 from sklearn.pipeline import Pipeline
11 from sklearn.preprocessing import StandardScaler, label_binarize
12 from sklearn.metrics import (
13     classification_report, confusion_matrix,
14     cohen_kappa_score, roc_auc_score,
15     f1_score, ConfusionMatrixDisplay
16 )
17
18 # -----
19 # 1. 4-class dataset: imbalanced (mirrors real document
20 # categorisation)
21 # -----
22 X, y = make_classification(
23     n_samples=3000, n_features=20, n_informative=12,
24     n_classes=4, n_clusters_per_class=1,
25     weights=[0.4, 0.3, 0.2, 0.1],    # class 3 is rare
26     random_state=42
27 )
28 class_names = ['Class A', 'Class B', 'Class C', 'Class D']
29 print(f"Class distribution: {np.bincount(y)}")
30
31 # -----
32 # 2. Model and out-of-fold predictions
33 # -----
34 cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
35 pipeline = Pipeline([
36     ('scaler', StandardScaler()),
37     ('clf', GradientBoostingClassifier(
38         n_estimators=200, max_depth=4,
39         learning_rate=0.05, random_state=42
40     )))
41
42 oof_labels = cross_val_predict(pipeline, X, y, cv=cv, n_jobs=-1)
43 oof_proba = cross_val_predict(pipeline, X, y, cv=cv,
44                               method='predict_proba', n_jobs=-1)
45
46 # -----
47 # 3. Full classification report (per-class precision/recall/F1)
48 # -----
49 print("\nClassification Report:")
50 print(classification_report(y, oof_labels, target_names=class_names))
51
52 # -----
53 # 4. Aggregate metrics with different averaging strategies
54 # -----
55 for avg in ['macro', 'weighted', 'micro']:
56     f1 = f1_score(y, oof_labels, average=avg)
57     print(f"F1 ({avg}):\t{f1:.4f}")
58
59 kappa = cohen_kappa_score(y, oof_labels)
60 print(f"Cohen's Kappa : {kappa:.4f} ")

```

```

61         f"({'good' if kappa > 0.6 else 'moderate' if kappa > 0.4 else
62             'poor'})")
63
64 # -----
65 # 5. Multiclass ROC-AUC (One-vs-Rest)
66 #     For multiclass AUC, we treat each class as binary (one vs. rest
67 #     )
68 #     and average. macro averages class AUCs equally; weighted by
69 #     support.
70 #
71 # -----
72
73 # 6. Normalised confusion matrix
74 #     Normalising by true labels (rows) shows recall per class.
75 #     Normalising by predicted labels (cols) shows precision per
76 #     class.
77 #
78 fig, axes = plt.subplots(1, 2, figsize=(13, 5))
79
80 cm_raw = confusion_matrix(y, oof_labels)
81 cm_norm = confusion_matrix(y, oof_labels, normalize='true')
82
83 for ax, cm, title, fmt in zip(
84     axes,
85     [cm_raw, cm_norm],
86     ['Confusion Matrix (counts)', 'Confusion Matrix (recall-
87         normalised)'],
88     ['d', '.2f']
89 ):
90     disp = ConfusionMatrixDisplay(confusion_matrix=cm,
91                                     display_labels=class_names)
92     disp.plot(ax=ax, colorbar=False, cmap='Blues', values_format=fmt
93     )
94     ax.set_title(title, fontsize=12, fontweight='bold')
95     ax.tick_params(axis='x', rotation=20)
96
97 plt.tight_layout()
98 plt.savefig("multiclass_confusion.png", dpi=150, bbox_inches='tight',
99 )
100
101 # -----
102 # 7. Per-class metrics visualisation
103 #
104 report = classification_report(y, oof_labels,
105                                 target_names=class_names, output_
106                                 dict=True)
107
108 metrics_df = pd.DataFrame({
109     'cls': {
110         'Precision': report[cls]['precision'],
111         'Recall': report[cls]['recall'],
112         'F1': report[cls]['f1-score'],
113     }
114 })

```

```

109         'Support' : report[cls]['support'],
110     }
111     for cls in class_names
112 ).T
113
114 fig, ax = plt.subplots(figsize=(9, 5))
115 x = np.arange(len(class_names))
116 w = 0.25
117
118 ax.bar(x - w, metrics_df['Precision'], w, label='Precision',
119         color='#0d47a1', alpha=0.85)
120 ax.bar(x, metrics_df['Recall'], w, label='Recall',
121         color='#00838f', alpha=0.85)
122 ax.bar(x + w, metrics_df['F1'], w, label='F1 Score',
123         color='#e65100', alpha=0.85)
124
125 ax.set_xticks(x)
126 ax.set_xticklabels(
127     [f'{c}\n{n=int(metrics_df.loc[c,'Support'])}' for c in class_
128         names],
129         fontsize=10
130 )
130 ax.set_xlim(0, 1.1)
131 ax.set_ylabel("Score", fontsize=11)
132 ax.set_title(f"Per-class Metrics (Kappa={kappa:.3f})",
133                 fontsize=12, fontweight='bold')
134 ax.legend(fontsize=10)
135 ax.grid(axis='y', alpha=0.3)
136 plt.tight_layout()
137 plt.savefig("multiclass_perclass.png", dpi=150, bbox_inches='tight')
138 plt.close()
139 print("\nMulticlass visualisations saved.")

140
141 # -----
142 # 8. Which averaging strategy to report?
143 # -----
144 print("\nAveraging Strategy Guide:")
145 print(f"  macro    F1={f1_score(y,oof_labels,average='macro'):.4f}
146      "
147      "-- equal weight to all classes (good for rare class
148          monitoring)")
149 print(f"  weighted F1={f1_score(y,oof_labels,average='weighted'):.4f}
150      "
151      "-- weight by support (good for overall performance)")
152 print(f"  micro    F1={f1_score(y,oof_labels,average='micro'):.4f}
153      "
154      "-- global TP/FP/FN (equals accuracy when all samples counted)
155      ")
156 print(f"  Kappa={kappa:.4f} -- corrected for chance (robust to
157          imbalance)")

```

13.4 The Normalised Confusion Matrix as a Debugging Tool

The raw confusion matrix shows counts; the normalised version (row-normalised by true class) shows recall per class. For a 4-class problem, reading down each column of the normalised matrix reveals which classes are being confused with each other. If Class

D (rare) is being predicted as Class A 60% of the time, the model has not learned a discriminative representation of Class D — a signal to collect more examples, oversample, or apply class-weighted loss.

Situation	Recommended Metric	Averaging
All classes equally important, balanced	Accuracy or Macro F1	Macro
Rare class matters most	Macro F1, per-class recall	Macro
Overall performance for a mixed audience	Weighted F1	Weighted
Comparing to human annotators	Cohen's Kappa	N/A
Probabilistic ranking across classes	Multiclass (OVR)	AUC

Table 6: Metric selection guide for multiclass problems.

Regression Models

- Reported MAE and RMSE (sensitivity to outliers differs).
- R^2 reported alongside residual plots — not in isolation.
- Four-panel residual diagnostic completed (patterns indicate missing structure).
- MAPE or sMAPE used only when target values are not near zero.

Calibration

- If probabilities are used directly (not just for ranking): reliability diagram inspected.
- Brier score reported alongside AUC.
- Calibration applied (Platt or isotonic) if reliability curve deviates from diagonal.

Statistical Testing

- Model comparisons confirmed with corrected resampled t-test or McNemar's test.
- Bonferroni correction applied when comparing more than two models.
- Difference in CV means reported with confidence interval, not just point estimate.

Fairness & Bias

- Protected attributes identified and excluded from model training.

- Per-group metrics computed for all relevant protected groups.
- Disparate Impact Ratio checked (≥ 0.8 required in most jurisdictions).
- Chosen fairness criterion (demographic parity vs equalized odds) documented and justified.

Production Monitoring

- PSI monitoring pipeline built before deployment, not after.
- Score distribution PSI and feature PSI monitored separately.
- Retraining triggers defined (e.g., $\text{PSI} \geq 0.2$ or AUC degradation $> 5\%$).
- Baseline reference distributions saved at deployment time.

Bias-Variance and Learning Curves

- Learning curves generated and inspected before finalising model complexity.
- Validation curve run for at least one key hyperparameter.
- Diagnosis (high bias / high variance / good fit) documented and acted upon.

Multiclass

- Averaging strategy (macro/weighted/micro) chosen based on class balance and business importance.
- Per-class precision, recall, and F1 inspected — not just aggregate metrics.
- Normalised confusion matrix inspected for inter-class confusions.
- Cohen's Kappa reported when comparing to human baseline or annotator agreement.

Appendix: Quick Reference

Key Formulas

$$\text{Precision} = \frac{TP}{TP + FP} \quad (1)$$

$$\text{Recall} = \frac{TP}{TP + FN} \quad (2)$$

$$F_\beta = (1 + \beta^2) \cdot \frac{P \cdot R}{\beta^2 P + R} \quad (3)$$

$$\text{Balanced Accuracy} = \frac{1}{2} \left(\frac{TP}{TP + FN} + \frac{TN}{TN + FP} \right) \quad (4)$$

$$\text{SHAP: } f(\mathbf{x}) = \mathbb{E}[f(X)] + \sum_j \phi_j \quad (5)$$

$$\text{PDP: } \hat{f}_j(x_j) = \frac{1}{n} \sum_{i=1}^n \hat{f}(x_j, \mathbf{x}_{-j}^{(i)}) \quad (6)$$

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2} \quad (7)$$

$$\text{Brier Score} = \frac{1}{n} \sum_{i=1}^n (\hat{p}_i - y_i)^2 \quad (8)$$

$$\text{PSI} = \sum_b (q_b - p_b) \ln(q_b/p_b) \quad (9)$$

$$\kappa = \frac{p_o - p_e}{1 - p_e} \quad (10)$$

Library Quick-Install

Required packages

```
pip install scikit-learn shap lime scikit-optimize scipy matplotlib
      pandas numpy seaborn
```

Key Classes and Functions

Class / Function	Purpose
StratifiedKFold	K-fold with class distribution preserved
TimeSeriesSplit	Walk-forward CV for temporal data
cross_validate	Multi-metric CV with train/test scores
cross_val_predict	Out-of-fold predictions for full-dataset evaluation
Pipeline	Chain preprocessing + model (prevents leakage)
GridSearchCV	Exhaustive hyperparameter search
RandomizedSearchCV	Stochastic hyperparameter search
BayesSearchCV	Bayesian hyperparameter optimisation (skopt)
average_precision_score	Area under precision-recall curve
balanced_accuracy_score	Accuracy corrected for class imbalance
precision_recall_curve	Threshold-free P-R analysis
shap.TreeExplainer	Fast SHAP for tree models
shap.summary_plot	Global SHAP visualisation
LimeTabularExplainer	Local LIME explanation for tabular data
PartialDependenceDisplay	PDP and ICE plots
mean_absolute_error	MAE for regression
r2_score	Coefficient of determination (R^2)
CalibratedClassifierCV	Platt scaling / isotonic regression
calibration_curve	Reliability diagram data
brier_score_loss	Probability calibration quality
learning_curve	Train/val score vs training set size
validation_curve	Train/val score vs hyperparameter value
cohen_kappa_score	Agreement metric for multiclass
roc_auc_score	AUC-ROC, supports multiclass OVR
ConfusionMatrixDisplay	Normalised multiclass confusion matrix

Acknowledgments



Esmaeil Rezaei, Ph.D.

Thank you for reading this guide! I hope this tutorial has helped you understand how to preprocess text data for transformer models. If you found this useful, I would greatly appreciate your support by connecting with me on my professional networks.

GitHub
[esmaeil-rezaei](#)

LinkedIn
[esmaeil-rezaei](#)

Feel free to explore my repositories on GitHub for more machine learning projects and connect with me on LinkedIn to stay updated on my latest work. Your feedback and suggestions are always welcome!